



EXECUTIVE GUIDE

The Executive's Guide to AI Fluency for Boards

How the human capability layer makes governance work.

By Zoë François

Founder, The AI Fluency Company

First edition · 2026

CONTENTS

What's in this guide.

FOREWORD	The third volume.	1
CHAPTER 1	The fluency–governance gap.	3
CHAPTER 2	The four competencies, briefly.	7
CHAPTER 3	Delegation at the executive level.	10
CHAPTER 4	Description as a leadership skill.	13
CHAPTER 5	Discernment and the limits of trust.	16
CHAPTER 6	Diligence and accountability.	19
CHAPTER 7	Mapping fluency to the ten pillars.	22
CHAPTER 8	Mapping fluency to the seven lifecycle stages.	25
CHAPTER 9	Building organisational fluency.	28
CLOSING	Three commitments.	30
ABOUT	The AI Fluency Company.	31

FOREWORD

The third volume.

Two guides told you what to govern. This one is about who is doing the governing — and whether they know how to work with AI well.

This is the third volume in a series of executive guides published by The AI Fluency Company. The first — The Executive's Guide to AI Governance — set out the ten pillars of an ISO/IEC 42001:2023-aligned AI Management System and the case for adopting one. The second — The Executive's Guide to the AI Product Lifecycle — took the lifecycle pillar and turned it into an operating manual. Together, they answer two questions: what should an organisation govern, and how should it govern it across the life of an AI system.

Those two guides assume something they do not deliver: that the humans doing the governing know how to work with AI in a competent, disciplined way. That assumption is increasingly unsafe. Most organisations now have AI tools in the hands of people who have never been taught how to use them well, supervised by leaders who have never been taught how to supervise them, governed by boards that have never been taught how to govern them. The policies are catching up. The people are not.

“AI policy without AI fluency is a building with rules but no occupants who can read them.”

The AI Fluency Framework, developed by Professor Joseph Feller (University College Cork) and Professor Rick Dakan (Ringling College of Art and Design) in partnership with Anthropic and published in 2025, gives the world a common vocabulary for the human capability layer. Four competencies — Delegation, Description, Discernment, Diligence — across three modes of interaction: Automation, Augmentation, Agency. The framework is published under a Creative Commons BY-NC-SA 4.0 licence and the original course is freely available at anthropic.com/ai-fluency and aifluencyframework.org. This guide does not reproduce it. There is a course for that, and you should take it.

What this guide does is apply it. It takes the four competencies and asks what they mean at the executive level, in a governed organisation, where the question is not 'how do I prompt this model better' but 'how do I lead an organisation in which thousands of decisions a day involve AI?' That is a different question. It needs a different treatment. And for South African boards, where external expectations are hardening and the local AI policy environment is still settling, it needs to be done now — not later, when those expectations have already become boardroom, customer, or assurance requirements.



A note on what this guide is not. It is not a prompt-engineering manual. It is not a tour of model capabilities. It is not a vendor guide. It is not a substitute for taking the underlying AI Fluency course — which any executive serious about this topic should do. It is, narrowly, a guide for the people whose job is to make organisational decisions about AI — set the strategy, allocate the budget, sign off the policy, brief the board, answer the customer, defend the decision when something goes wrong. Those people need fluency too. They need a different shape of fluency than a power user does. And they rarely get the training that other people get, because there is an unspoken assumption that they will pick it up by osmosis. They will not.

The structure of this guide reflects that. Chapter 1 makes the central argument: that there is a gap between AI governance and AI execution, and that fluency is the layer that closes it. Chapter 2 covers the framework itself, briefly and respectfully — just enough to ground the rest. Chapters 3 to 6 take each of the four competencies and translate them for the executive context, with the same structure as the lifecycle guide: what the competency requires, what good looks like, where organisations go wrong, three questions to ask. Chapters 7 and 8 connect fluency to the ten governance pillars and the seven lifecycle stages from the previous volumes. Chapter 9 gives a practical roadmap for building organisational fluency. The closing chapter restates the argument as three commitments.

The South African context, as ever, matters. Most boards I work with are still in the position where the policy is being drafted, the inventory is being built, the committee is being formed. That is the moment to build fluency in. Not after. The April 2026 withdrawal of South Africa's Draft National AI Policy — pulled after journalists discovered fabricated citations generated by AI used without adequate human oversight (Government Gazette Notice 3880 of 2026, gazetted 10 April, withdrawn 26 April) — is the sharpest possible illustration of what happens when governance exists but fluency does not. A government that could not govern its own AI-assisted document is a government that has learned, viscerally, why the human capability layer is not optional. Once the policy is signed and the committee meets quarterly, the chance to get this right has narrowed. The people now using AI in your organisation are forming habits this month that will define how AI works for you for the next five years. Fluency is what shapes those habits.

Zoë François

Cape Town, 2026

CHAPTER 1

The fluency–governance gap.

Two disciplines, both partial. The organisations that get AI right will run them together.

Most organisations are now investing in one of two responses to AI. Some are investing in governance — policies, committees, ISO/IEC 42001:2023 readiness, vendor reviews, risk assessments. Others are investing in capability — training, tools, prompt libraries, internal academies, pilot projects. Almost none are investing in both, in equal measure, with the same seriousness, owned by the same leadership team. That asymmetry is the source of most AI programme failures I see, and it is the gap this guide is written to close.

The two failure modes.

Organisations that invest only in governance produce a familiar pattern. The policy is signed. The committee meets. The risk register has entries. The audit programme is scheduled. And in the operational layer, where the work actually happens, AI is being used in ways the policy doesn't cover, by people who haven't read it, in tools the committee doesn't know are deployed. The governance is real. The work is unchanged. The gap between the two is where incidents happen.

Organisations that invest only in capability produce a different and equally familiar pattern. The training is rolled out. People are excited. Productivity stories accumulate. Use cases multiply. And then someone deploys AI into a customer-facing process without a privacy assessment, or a manager uses an AI summary in a performance decision, or a sensitive document is uploaded to a tool that is not contractually approved. The capability is real. The discipline is missing. The gap between the two is where reputational damage happens.

Both failure modes share a common cause: the organisation has invested in one half of a two-part system. AI works in organisations only when governance and fluency are present together. Governance without fluency produces theatre — the documents exist, but the work is uncontrolled. Fluency without governance produces drift — the work is skilful, but the constraints are invented locally. The board's job is to ensure both are funded, both are owned, and both are built at the same rate.

“Governance answers what should be controlled. Fluency answers how a competent human exercises that control. Neither works alone.”

What ISO/IEC 42001 assumes about people.

Read ISO/IEC 42001:2023 closely and you will find a quiet assumption running through it: that the humans inside the management system are competent to do what the system asks of them. The standard requires human oversight of high-stakes AI decisions. It does not say what makes that oversight meaningful. The standard requires risk assessment for AI use cases. It does not say how to recognise an AI risk when you see one. The standard requires transparency with stakeholders about AI use. It does not say how to communicate it well. In every case, the standard prescribes the requirement and leaves the capability question open.

That is by design. ISO standards are deliberately capability-agnostic — they tell you what to do, not whether your people can do it. The implicit assumption is that the management system will be operated by competent professionals. For a quality management system, that assumption is reasonable. The professions involved — quality engineers, auditors, operations leaders — have been trained for decades. For an AI management system in 2026, the assumption is unsupported. Most of the people now operating AI tools in mid-market organisations have never been trained on AI at all. The standard expects them to exercise judgement they have not yet been taught.

This is what I mean by the fluency–governance gap. The standard is sound. The framework is settled. The pillars are clear. But the human layer the standard relies on is, in most organisations, undertrained, uncertified, and operating on instinct. Fluency is the discipline of building that layer. It is not a substitute for governance. It is what governance is silently assuming, and what most organisations have not yet provided.

What changes when fluency is present.

The change is not that incidents disappear — they don't. The change is that the organisation starts catching things it would otherwise have missed. A fluent employee notices when a tool is producing confidently wrong outputs. A fluent manager recognises when an AI-assisted decision is being rubber-stamped rather than reviewed. A fluent procurement leader asks vendors questions that surface model risk. A fluent executive can read an AI vendor's pitch and tell the difference between substance and theatre. None of this requires deep technical knowledge. It requires the four competencies of the AI Fluency Framework, applied at the level of judgement appropriate to the role.


Fluency also changes how the governance system itself operates. A risk register written by people who don't understand AI risk lists the wrong risks. An audit conducted by people who can't recognise model drift will not find it. A board briefing produced by people who can't evaluate AI evidence will accept the version the vendor wants them to accept. Governance is only as good as the people inside it, and AI is the first technology in a generation where the people inside the governance system have to learn the technology before the governance can be trusted.

The remaining chapters of this guide do two things. They translate the four AI Fluency competencies into something a senior leader can use. And they connect those competencies to the ten governance pillars and seven lifecycle stages from the earlier volumes, so that the integration of fluency and governance becomes a single, coherent discipline rather than two parallel programmes.

THE BOARD TEST

- Could you, today, name the AI tools your largest department uses every week?
- If a major customer asked how your organisation decides what AI is appropriate for which task, who would answer, and what would they say?
- When was the last time the executive team had a substantive conversation about how AI is being used by your employees — not whether to invest in it, but how it is being used, today?

If those questions are uncomfortable, the gap is real. The remaining chapters are about closing it.



CHAPTER 2

The four competencies, briefly.

A respectful summary of someone else's framework — just enough to ground the rest of this guide.

This chapter is short by design. The AI Fluency Framework is the work of Professor Joseph Feller (University College Cork) and Professor Rick Dakan (Ringling College of Art and Design), in partnership with Anthropic. It was published in 2025 under a Creative Commons BY-NC-SA 4.0 licence. The original course — AI Fluency: Framework and Foundations — is freely available and has attracted more than 300,000 enrolments globally (HEA, April 2026). It runs to several hours of video and structured exercises. If you intend to lead an organisation's use of AI, you should take the course. There is no substitute, and this chapter does not pretend to be one.

What follows is the minimum needed to read the rest of this guide. The framework defines four core competencies necessary to interact with AI in ways that are effective, efficient, ethical, and safe. It also identifies three modes of AI interaction that the competencies apply across. Both are summarised below in the framework's own terms.

The four competencies (the 4Ds).

DELEGATION

- Setting goals and deciding whether, when, and how to engage with AI.
- Includes: understanding the problem you're solving (problem awareness), understanding what AI systems can and cannot do (platform awareness), and distributing tasks thoughtfully between humans and AI (task delegation).

DESCRIPTION

- Effectively communicating with AI systems to prompt useful behaviours and outputs.
- Includes: providing context, examples, and constraints; specifying tone and role; encouraging step-by-step reasoning; and structuring the request so the AI can produce something useful on the first attempt.

DISCERNMENT

- Accurately assessing the usefulness of AI outputs, processes, and behaviours.
- Includes: evaluating products (is the output correct, complete, appropriate?), processes (is the AI reasoning soundly?), and performance (is the system behaving consistently over time?).

DILIGENCE

- Taking responsibility for what we do with AI and how we do it.
- Includes: creation diligence (the choice of tools and data), transparency diligence (disclosing AI involvement), and deployment diligence (vouching for the final product, accepting responsibility for its impact).

The framework also describes a Description–Discernment loop: a working pattern in which a user describes a goal, evaluates the AI’s response, refines the description, and iterates until the output is fit for purpose. This loop is the operating mechanic of most augmented work, and competent users move through it dozens of times a day, often without realising it.

The three modes of AI interaction.

THREE MODES

Automation

- AI executes specific tasks based on human instructions. The human defines the work; the AI does it.

Augmentation

- Humans and AI collaborate as thinking partners — brainstorming, drafting, reviewing, deciding together.

Agency

- Humans configure AI to perform future tasks independently on their behalf. The human shapes the AI’s knowledge and behaviour rather than instructing it task by task.

Each of the four competencies applies across all three modes — but they show up differently in each. Delegation in automation looks like task selection. Delegation in agency looks like architecture. Discernment in augmentation looks like real-time evaluation of a draft. Discernment in agency looks like setting up automated checks the AI cannot bypass. The shape of fluency depends on the mode of use.

Two further observations matter for the rest of this guide. First, the four competencies are interconnected, not sequential. A weakness in one cascades into the others. Poor delegation — asking AI to do a task it cannot do well — leads to wasted description, sceptical discernment, and diligence that has to clean up the mess. Strong delegation makes the other three easier. Second, the framework is platform-agnostic. It is not about ChatGPT or Claude or Copilot. It is about the human capabilities required to work with any AI system effectively, efficiently, ethically, and safely. That generality is what makes it useful for governance: the framework outlives any particular vendor.

“The four competencies are not a checklist. They are a way of thinking about whether the humans in your organisation can be trusted with the AI you have already given them.”

With that grounding in place, the rest of this guide takes each competency in turn and asks the question this guide is uniquely able to ask: what does each competency look like at the executive level, in a governed organisation, with a board to brief and a regulator to satisfy?

A NOTE ON ATTRIBUTION

The AI Fluency Framework is the intellectual property of Rick Dakan, Joseph Feller, and Anthropic, published in 2025 under the Creative Commons BY-NC-SA 4.0 licence. This guide applies the framework to the executive context. The framework descriptions in this chapter are summaries; for the authoritative source, take the free Anthropic course at anthropic.com/ai-fluency or visit aifluencyframework.org.

CHAPTER 3

Delegation at the executive level.

Boards don't prompt AI. They decide what AI is appropriate or permitted to do, if at all. That is the executive form of delegation — and most boards are doing it badly.

Delegation is the most consequential of the four competencies for executive teams, and the one most often misunderstood. At the user level, delegation is the choice of which task to give to AI. At the executive level, delegation is the choice of which categories of work the organisation will allow AI to do, in what conditions, with what oversight, with what accountability. It is the meta-decision about AI — the one all the operational decisions depend on.

Most boards have not made this decision. They have allowed it to be made for them, by hundreds of small purchasing decisions, vendor pitches, and quietly-rolled-out features in tools they already pay for. The result is an organisation in which AI is doing many things, often well, sometimes badly, and the board has no view of which is which. Executive-level delegation is the discipline of taking that decision back.

What this competency covers at the executive level.

The AI Fluency Framework defines delegation as setting goals and deciding whether, when, and how to engage with AI — with three sub-competencies: problem awareness, platform awareness, and task delegation. At the executive level, each of these scales up.

Problem awareness becomes strategic awareness: which of the organisation's actual business problems are AI-shaped, and which only look like they are because someone is selling a tool?

Platform awareness becomes capability awareness: which classes of AI system can your organisation responsibly operate, given its data, its skills, its risk appetite, its sector? Task delegation becomes domain delegation: which categories of work is AI permitted to do here, in what conditions, with what oversight, and which is it not?

These are governance decisions that look like capability decisions. They cannot be made by the technology team alone. They cannot be made by the compliance team alone. They have to be made by an executive group that understands both halves — what is technically possible, and what is organisationally appropriate. That group does not, in most organisations, exist yet. Executive-level delegation is the discipline of forming it.

WHAT STRONG EXECUTIVE DELEGATION REQUIRES

- A documented organisational stance on which categories of work AI is permitted to do — and which are off the table.
- A short list of high-trust use cases where AI delivers measurable value, owned by named senior leaders.
- A small set of explicit prohibitions — use cases the organisation has decided not to pursue, with reasons documented.
- A policy that distinguishes between AI use that can proceed under normal management authority and AI use that requires executive sign-off.
- An understanding, at executive level, of the difference between automation, augmentation, and agency — and which mode each major AI use case is operating in.

WHAT GOOD LOOKS LIKE

- The board can name the five highest-value AI use cases in the organisation, who owns each, and what success looks like.
- The board can name three categories of AI use that have been deliberately ruled out — and explain why.
- Procurement decisions involving AI flow through a documented delegation framework rather than through individual managers' budgets.
- New AI use cases enter through a defined intake process with a named senior sponsor and a documented business case.
- There is a clear distinction between 'AI we are using as a tool' and 'AI we are configuring to act on our behalf' (agency), with the latter receiving more rigorous review.
- The organisation's risk appetite for AI is documented in language that operational managers can apply.

WHERE ORGANISATIONS GO WRONG

- Letting delegation happen by accident — every team chooses its own tools, the board sees only the aggregate.
- Treating AI as an IT decision rather than a strategic one — letting the technology team scope what AI is for.
- Saying yes to AI tools without saying no to anything — a permissive policy with no explicit prohibitions is not a policy.
- Confusing automation with agency — approving an 'AI tool' without realising it has been configured to take actions autonomously.
- Outsourcing the delegation question to vendors — letting the vendor decide what good looks like for your business.

- Failing to revisit delegation as the technology evolves — a stance set in 2024 is already out of date.

THREE QUESTIONS TO ASK YOUR ORGANISATION

- What categories of work has our organisation explicitly decided AI will and will not do — and where is that decision documented?
- Of the AI tools currently in use across the business, how many entered through a deliberate executive-level delegation decision, and how many entered through other routes?
- Which of our current AI uses are operating in agency mode — making decisions or taking actions autonomously — and who approved that?

The mode question — and why it is a board question.

The three modes of AI interaction — automation, augmentation, agency — sound like a technical distinction. They are a governance distinction. Automation, where AI executes specific tasks the human has defined, has a relatively narrow risk surface and can usually be governed at operational level. Augmentation, where humans and AI collaborate on judgement work, has a broader risk surface and needs management-level oversight. Agency, where AI is configured to act on the organisation's behalf in ways the human will not review case by case, has the broadest risk surface of any technology category in modern business — and needs board-level visibility.

Most boards I work with have approved agentic AI without realising they have done so. A vendor proposes an 'AI assistant' for the customer service team. The contract is signed at functional level. It turns out the assistant can resolve cases, issue refunds, escalate to managers, and update customer records — all without human review. That is agency, sold as automation. The board did not approve it because the board was not asked. The first the board hears about it is when something goes wrong.

The discipline of executive-level delegation is, above all, the discipline of recognising the mode of each material AI use case and ensuring the level of oversight matches. Agency requires more than automation requires. The board's job is to ensure that distinction is being made — and that the boundary between what AI can do without approval and what it cannot is drawn, and known, and respected.

THE MODE QUESTION — WHY IT IS A BOARD QUESTION

The three modes carry very different risk profiles. **Automation** (AI executes specific tasks the human defines) has a relatively narrow risk surface and can usually be governed at operational level. **Augmentation** (humans and AI collaborate on judgement work) has a broader risk surface and needs management-level oversight. **Agency** (AI is configured to act on the organisation's behalf without case-by-case human review) has the broadest risk surface of any technology category in modern business — and requires board-level visibility.

Most boards have approved agentic AI without realising they have done so. A vendor proposes an AI assistant. The contract is signed at functional level. The assistant turns out to be resolving cases, issuing refunds, and updating records — all without human review. That is agency, sold as automation. The board's job is to ask: which of our AI systems are operating in agency mode, and who approved that?

“The most expensive AI delegations an organisation makes are the ones it makes by accident.”

CHAPTER 4

Description as a leadership skill.

Leaders don't write prompts. They set context. The same discipline applies, at a different level.

Description is the competency most often reduced to a slogan: prompt engineering. At the user level, prompt engineering is real and useful. At the executive level, the relevant skill is broader and more important. Leaders set the context within which thousands of AI interactions happen across the organisation every day. The clarity of that context — what the organisation is for, what it values, what it considers acceptable, what it considers off-limits — shapes the quality of every prompt every employee writes. Most organisations have never thought of leadership as a description discipline. It is.

This chapter takes the AI Fluency Framework's second competency — effectively communicating with AI systems to prompt useful behaviours and outputs — and applies it at the level where executives operate. Not how to phrase a single prompt. How to communicate, across the organisation, the framing within which good prompts can be written by everyone else.

What description means at the executive level.

Three things change when description scales from individual to organisational. First, the audience changes: leaders are not communicating with the AI directly, they are communicating with the humans whose communications with AI shape the organisation's output. Second, the scope changes: a single prompt is concerned with one task, but an executive description is concerned with the standing context within which all tasks happen. Third, the stakes change: a poor individual prompt produces a poor individual output. A poor organisational description produces poor outputs across thousands of employees, often invisibly, often with a lasting effect.

Concretely, executive description shows up in four artefacts. The AI policy — which is a description of the boundaries within which AI can be used. The internal communications about AI — which describe the organisation's stance, expectations, and limits to employees. The vendor briefs — which describe what the organisation needs from an AI partner. The customer-facing communications — which describe how AI is being used in the service the organisation delivers. Each of these is a description, in the AI Fluency Framework's sense. Each shapes how AI behaves in the organisation. Each is usually written badly.

WHAT STRONG EXECUTIVE DESCRIPTION REQUIRES

- An AI policy written in plain language, with concrete examples, that an operational manager can apply without needing legal interpretation.
- Standing organisational context that any employee can refer to when working with AI — what the organisation is for, what it values, what it forbids.
- Clear role-based guidance — the description of how AI should be used in finance differs from how it should be used in marketing, and both differ from customer service.
- Vendor briefs that describe outcomes the organisation wants, not the technologies it has heard about.
- Customer-facing language that is honest about where AI is in the service — written in customer language, not legal disclaimers.
- Internal communications that update this context as the organisation's stance evolves — description is not a one-time act.

WHAT GOOD LOOKS LIKE

- An employee can answer the question 'is it OK to use AI for this?' without escalating, in most cases, by referring to a single document.
- AI use cases produced by different teams have a recognisable consistency — they sound like outputs from the same organisation.
- The AI policy is short, written for non-specialists, and revised at least annually as the technology and the organisation evolve.
- Customer disclosures about AI are clear, accurate, and proportionate — not hidden in terms of service, not exaggerated in marketing.
- Vendor pitches are evaluated against a written brief that describes the outcome wanted, not the feature set.
- The executive team can articulate, in a single paragraph, what role AI plays in the organisation's strategy.

WHERE ORGANISATIONS GO WRONG

- Confusing prompt engineering for executive description — thinking the skill is at user level, missing it at organisational level.
- Writing AI policies that no operational manager can apply — documents that are technically complete but practically useless.
- Communicating about AI only when something goes wrong — leaving the standing context to be filled in by rumour.
- Letting customer-facing communication about AI default to legal language — producing disclosures that obscure rather than inform.
- Treating description as a one-time act — writing the policy in 2025 and not revisiting it.

- Letting different parts of the organisation produce divergent descriptions of what AI is for — marketing says one thing, HR says another, the policy says a third.

THREE QUESTIONS TO ASK YOUR ORGANISATION

- If we asked a random employee tomorrow ‘what is our organisation’s position on using AI for this kind of work?’, how consistent would the answers be?
- When was our AI policy last revised, and could a new manager apply it without escalating to legal or compliance?
- How do our customers learn about where AI is used in the service we provide — and would they describe what we tell them as clear or as evasive?

The Description–Discernment loop, scaled up.

The framework describes a working loop: describe to the AI, evaluate the response, refine the description, evaluate again. At the organisational level, the same loop operates over months rather than minutes. The organisation describes its AI stance — in policy, in communications, in vendor briefs. It then sees how that stance is interpreted in practice — by employees, by vendors, by customers. It evaluates the gap between intention and interpretation. And it refines its description: the policy is updated, the communications are clarified, the briefs are sharpened.

Most organisations do not run this loop. They write the policy once, communicate it once, and then treat any deviation as a failure of compliance rather than as data about the description being unclear. The mature organisation reads deviations as feedback. If twenty employees are interpreting the policy in the same wrong way, the policy is the problem, not the employees. That insight is the executive form of discernment, applied to the executive form of description. The next chapter deals with discernment proper.

“The clarity of leadership shapes the clarity of every prompt every employee writes. Most leaders have never thought of themselves as the prompt.”

Discernment and the limits of trust.

AI produces confidence-shaped outputs. The discipline of telling whether it is right is the rarest of the four — and the one that protects the organisation.

Discernment is the competency that sounds easiest and is hardest. The AI Fluency Framework defines it as accurately assessing the usefulness of AI outputs, processes, and behaviours. At user level, that is the work of fact-checking, sense-checking, and rejecting AI output that doesn't hold up. At executive level, it is the work of building an organisation in which discernment is structured, expected, and resourced — because relying on individual users to exercise it case by case is, in most operational contexts, an inadequate control.

The reason discernment is hard is the same reason AI is useful: the outputs are confident, fluent, and superficially competent. A model that produced obviously wrong answers would be easy to discern. A model that produces plausibly wrong answers — well-structured, professionally worded, occasionally fabricated — is a different problem. It defeats casual review. It rewards the user who wants to ship and punishes the user who wants to verify. It encodes the failure mode of every modern AI deployment. Discernment is the discipline of resisting it.

What discernment means at the executive level.

The AI Fluency Framework distinguishes three forms of discernment: of the product (the output the AI produces), of the process (the reasoning the AI follows), and of the performance (how the system behaves over time). Each scales differently to the executive level. Product discernment scales into quality assurance — how does the organisation systematically check that AI outputs are fit for purpose? Process discernment scales into auditability — can the organisation explain how AI decisions are being made? Performance discernment scales into monitoring — does the organisation know when an AI system is degrading?

All three are governance questions. None can be answered by training individual users to be sceptical — although that helps. They require organisational structures: review processes, audit programmes, monitoring systems, escalation paths. Discernment is where AI fluency meets ISO/IEC 42001:2023 most directly. The standard's requirements for performance evaluation, internal audit, and continual improvement are, in effect, discernment requirements applied at the system level.

WHAT STRONG EXECUTIVE DISCERNMENT REQUIRES

- Defined quality criteria for AI outputs in each material use case — what 'good' looks like, in measurable terms.
- A review discipline proportionate to the stakes — high-impact AI outputs receive more scrutiny than low-impact ones.
- Process auditability — the organisation can explain how an AI decision was reached, sufficient to satisfy a stakeholder enquiry.
- Monitoring of AI performance over time, with thresholds that trigger investigation when crossed.
- Independent review where appropriate — high-stakes AI use cases evaluated by people who did not commission them.
- A culture in which questioning AI output is rewarded, not penalised — most organisations get this culture wrong without realising.

WHAT GOOD LOOKS LIKE

- AI outputs in customer-facing or high-stakes contexts pass through a defined review step before they are issued.
- There is a documented standard for what constitutes adequate review — not a generic 'check it before you send it'.
- When AI output is questioned, the question is treated as data about the system, not as friction in the workflow.
- Internal audit treats AI systems as in scope, with auditors who have been trained to evaluate them — not avoided as too technical.
- Performance metrics for AI systems are reported to executive level as part of operational reporting, not buried in technical dashboards.
- Drift, bias shifts, and unusual outputs are detected by the system's monitoring rather than by an aggrieved stakeholder.

WHERE ORGANISATIONS GO WRONG

- Treating AI fluency as confidence — rewarding users who use AI prolifically rather than users who use it carefully.
- Building AI into workflows in ways that make discernment expensive — forcing employees to choose between checking and shipping.
- Outsourcing discernment to the AI itself — 'the model said it had checked' is not a control.
- No defined quality standard for AI output — leaving every reviewer to invent their own.
- Punishing employees who flag AI errors — producing a culture in which mistakes are quietly absorbed rather than systemically captured.

- No monitoring of AI performance over time — detecting drift only when it produces a visible incident.
- Internal audit programmes that exclude AI on grounds of technical complexity — leaving the organisation's most consequential systems uninspected.

THREE QUESTIONS TO ASK YOUR ORGANISATION

- For our highest-stakes AI use case, what is the defined review step before output is issued — and could you describe it?
- When was the last time an AI error was detected, reported, and fed back into improvement — and what changed as a result?
- If our AI systems started drifting tomorrow, who would notice, what would they notice it on, and how long would it take?

The discernment culture problem.

Most discernment failures are not technical. They are cultural. An organisation that has trained its employees to use AI but not to question it produces employees who use AI confidently and review it thinly. An organisation that measures productivity gains from AI without measuring quality loss produces a metric that looks good and an output that gets worse. An organisation that frames AI scepticism as resistance to change produces a workforce that has learned to keep its scepticism quiet. All three patterns are common. All three are the executive's problem to fix.

The fix is not difficult to describe but it is unfashionable. It requires the executive team to model discernment publicly: to ask, in board meetings, where the AI evidence comes from. To send a draft back when the source is opaque. To compliment employees who flag AI errors and to thank them for catching it. To resource review steps in workflows even when it slows things down. To measure AI quality alongside AI productivity. To normalise the question 'how do we know this is right?' as something one asks of every AI-derived input, including ones from senior people who appear confident.

This is not an anti-AI stance. It is a fluent one. Fluent users are not afraid of AI; they trust it within the boundaries within which it has earned trust, and they verify outside those boundaries. That posture has to be modelled from the top. If the executive team accepts AI evidence uncritically, no amount of training elsewhere in the organisation will produce employees who don't.

The Description–Discernment loop in operational reality.

The loop the framework describes — describe, evaluate, refine, repeat — is what competent users do dozens of times a day. It is also what mature organisations do, structurally, around their AI systems. Each operational AI system should run a version of this loop: the system's outputs are described in monitoring metrics, evaluated against thresholds, refined when drift is detected,

evaluated again. The loop is the operating mechanic of continuous improvement, which is the eighth pillar of governance.

Where most organisations fail is on the second half of the loop. They describe their AI systems carefully at deployment. They then do not evaluate them in operation, do not refine them in maintenance, do not close the loop. The Description–Discernment loop is the test of whether an organisation's AI is being managed or merely deployed.

“AI does not produce fewer errors. It produces more confident ones. Discernment is the discipline that prevents confidence from becoming evidence.”

CHAPTER 6

Diligence and accountability.

Of the four competencies, this is the one that sits closest to governance — and the one that makes governance real.

Diligence is the competency the AI Fluency Framework defines as taking responsibility for what we do with AI and how we do it. It is the ethics-and-safety dimension of fluency — distinct from the effectiveness dimensions of delegation, description, and discernment. The framework distinguishes three forms: creation diligence (the choice of tools and data, with awareness of their privacy, security, and ethical track record), transparency diligence (disclosing AI involvement to relevant stakeholders), and deployment diligence (vouching for the final product, accepting responsibility for its impact).

At the executive level, diligence is what most directly maps onto the governance system the previous two volumes described. The ten pillars are, in effect, an organisational architecture for diligence. ISO/IEC 42001:2023's requirements for accountability, transparency, risk management, and stakeholder trust are all forms of diligence at the system level. This chapter takes the framework's definition of diligence and asks what it requires of an executive team — not as a virtue, but as a discipline.

What diligence means at the executive level.

The clearest way to understand executive diligence is through the question of who answers when something goes wrong. In organisations with weak diligence, that question has no clean answer. The user who deployed the AI says they followed policy. The vendor says they delivered the tool as specified. The compliance team says they reviewed the policy, not the deployment. The board says they approved a strategy, not a use case. Everyone is partly right and nobody is accountable. That is the failure mode diligence prevents.

In organisations with strong diligence, accountability is named in advance, not negotiated after the fact. Each material AI use case has a named senior owner. Each high-stakes decision has a documented sign-off. Each customer-affecting AI system has a documented disclosure. Each AI system has a documented owner of its outputs. None of this is bureaucratic in the bad sense; all of it is the minimum required to be able to answer, in real time, the question every regulator and every customer eventually asks: who is responsible for this?

WHAT STRONG EXECUTIVE DILIGENCE REQUIRES

- Named accountability for every material AI use case — not a department, a person.
- Documented choice of tools, data, and vendors — with the rationale preserved, not just the decision.
- Disclosure mechanisms for AI involvement, calibrated to the audience: customers, employees, regulators, partners.
- An organisational stance on attribution — when an AI-assisted output goes out, who is vouching for it, and on what basis.
- Incident response procedures specific to AI — what to communicate, to whom, in what timeframe, when something goes wrong.
- Standing acceptance of responsibility — the organisation does not blame its tools for outcomes, even when the tool was the proximate cause.

WHAT GOOD LOOKS LIKE

- If an AI system produced a bad outcome tomorrow, the organisation could name within minutes who is accountable, what disclosures have been made, and what the response will be.
- Customer-facing communications about AI are honest: they describe what the AI does, where the human is in the loop, and how to raise concerns.
- Employees know how AI affects their work — their performance reviews, their data, their daily tools — and have a channel to ask questions.
- Vendor due diligence covers the AI components specifically, including how the vendor uses customer data and what happens to that data when the contract ends.
- There is a documented incident response plan for AI failures, tested at least once.
- When the organisation is asked 'who is responsible?', the answer arrives quickly and is the same regardless of who is asking.

WHERE ORGANISATIONS GO WRONG

- Treating diligence as compliance — producing the documents required for an audit but not the discipline they describe.
- Allowing accountability for AI to land on whoever happens to be using it — a fragmented arrangement that holds up only as long as nothing goes wrong.
- Concealing AI involvement from customers — either deliberately, or through legal language designed to obscure.
- Treating AI errors as user errors — punishing the employee who deployed the tool rather than improving the system.
- Outsourcing diligence to vendors — trusting that the AI provider has handled ethics, privacy, and safety on the organisation's behalf.
- Having no incident response plan for AI — inventing one in real time when an incident occurs.

- Confusing the absence of an incident with the presence of diligence — not yet caught is not the same as not happening.

THREE QUESTIONS TO ASK YOUR ORGANISATION

- If our most consequential AI system produced harm tomorrow, who would answer for it — by name, before the meeting starts?
- How do our customers and employees learn about where AI is in our processes, and would we be comfortable defending the disclosure in front of a regulator?
- When did we last test our AI incident response plan — and if we don't have one, why not?

The three forms of diligence at the organisational level.

The AI Fluency Framework's three forms of diligence — creation, transparency, and deployment — each take on an organisational shape that is worth naming explicitly.

Creation diligence at the organisational level is the discipline of vendor and tool selection. It includes the choice of which AI systems to bring into the organisation, on what contractual terms, with what data flows, with what alignment to the organisation's stated values. Most organisations do this badly because the procurement function does not yet know what to ask of AI vendors. The questions that matter — about data residency, model provenance, handling of customer data after contract end, ability to comply with subject access requests — are not the questions standard procurement asks. Creation diligence requires upgrading procurement's capability deliberately.

Transparency diligence at the organisational level is the discipline of disclosure. It includes how the organisation tells customers AI is in use, how it tells employees their work involves AI, how it tells regulators what AI is doing, how it tells investors and auditors what AI risk it is carrying. Each audience needs different language, different detail, different timing. Most organisations have no coherent disclosure strategy for AI — they have whatever each function happens to publish, often inconsistent with what other functions publish. Transparency diligence is the discipline of making those disclosures coherent and accurate.

Deployment diligence at the organisational level is the discipline of accepting responsibility for AI-mediated outcomes. The organisation's name is on the output, regardless of how the output was produced. 'The AI did it' is not a defence available to the organisation — the organisation chose to deploy the AI, configured it, gave it access, monitored it (or didn't). Deployment diligence is the discipline of standing behind AI outputs the way the organisation stands behind any other output it produces. Most organisations have not yet decided whether they will. The decision needs to be explicit, and made before an incident, not during one.

“Diligence is what turns governance from a document into a discipline. Without it, the policy is a story. With it, the policy is a system.”

The Diligence Statement — a small habit with large consequences.

The original AI Fluency course suggests the practice of writing a Diligence Statement: a short, honest disclosure that accompanies AI-assisted work, naming what AI was used, for what, and what the human is taking responsibility for. At individual level, it is a small habit. At organisational level, it is a culture-shaping practice. An organisation in which AI-assisted work is signed with a clear statement of human accountability is an organisation in which accountability stays where it should be. An organisation that suppresses such statements — because they look bureaucratic, or might unsettle stakeholders, or might invite scrutiny — is an organisation that has chosen not to know what it is doing.

The executive decision to encourage Diligence Statements (in some form, calibrated to context) is small in itself. The signal it sends is large. It says: AI is being used here, openly, with the human accountable for the output. That signal is what stakeholders are increasingly asking for. The organisations that get there first will benefit; the ones that resist will eventually be required to do it anyway, in worse circumstances.

CHAPTER 7

Mapping fluency to the ten pillars.

How the four competencies show up in the governance system the first volume described.

The ten pillars set out in The Executive's Guide to AI Governance organise the requirements of ISO/IEC 42001:2023 into a structure boards can use. Each pillar is a category of work the organisation has to do. The four AI Fluency competencies are the human capabilities that determine whether each pillar's work is done well or poorly. This chapter maps the two.

The mapping is not one-to-one. Each pillar draws on more than one competency, and each competency contributes to more than one pillar. What the mapping shows is the dominant relationship: the competency that, if absent, will most undermine the pillar. Read the table below as a diagnostic. If a particular pillar is weak in your organisation, the corresponding competency is the most likely cause.

PILLAR	DOMINANT FLUENCY COMPETENCY — AND WHY
1. Strategic Intent	Delegation. Strategic intent is the answer to 'what problem is AI solving for us?' — which is the highest-level form of delegation. Without delegation discipline at executive level, strategic intent collapses into vendor-driven aspiration.
2. Governance & Leadership	Diligence. Real accountability is a diligence outcome. The named owner, the documented decision rights, the willingness to stand behind AI outputs — all are organisational diligence in its executive form.
3. Risk & Impact Management	Discernment. Most AI risks are recognisable only to people fluent enough to see them. A risk register written by people without discernment lists the wrong risks — or, more dangerously, the obvious ones, missing the subtle ones.
4. AI Lifecycle Management	All four. Lifecycle management is the longitudinal application of fluency — delegation at concept, description at design, discernment at validation and operation, diligence at deployment and decommissioning.
5. Data & Technology Governance	Diligence (creation). What runs the AI — the data, the platforms, the third-party services — is the artefact of organisational creation diligence. Weak diligence here produces brittle, opaque, hard-to-govern AI estates.
6. Human Oversight	Discernment. Oversight is not the presence of a human in the workflow; it is a human exercising discernment in the workflow. Without trained discernment, oversight is theatre.

PILLAR	DOMINANT FLUENCY COMPETENCY — AND WHY
7. Performance Evaluation	Discernment. Performance evaluation is structured organisational discernment — measuring outputs, processes, and longitudinal behaviour. The Description–Discernment loop, run at system level.
8. Continual Improvement	Discernment and diligence together. Improvement requires discerning what to change and the diligence to actually change it — most organisations have one without the other.
9. Compliance & Assurance	Diligence (transparency and deployment). Assurance is the capacity to evidence, to a stakeholder, that the organisation’s diligence has been real — not described, but practised.
10. Stakeholder Trust & Transparency	Diligence (transparency). The cleanest mapping in the framework. Stakeholder trust is the organisational form of transparency diligence — honest, audience-appropriate, sustained over time.

Two patterns emerge from the mapping. First, every pillar is supported by at least one competency — there is no governance work that fluency does not enable. Second, two competencies appear repeatedly: discernment and diligence. Those are the competencies most often weak in organisations, and the ones most directly responsible for governance failures. An organisation that wants to strengthen its governance posture should start with those two.

USING THE MAPPING DIAGNOSTICALLY

- For each pillar where your governance is weak, identify which competency the mapping points to.
- Audit the relevant population for that competency — do the people responsible for that pillar have the fluency they need?
- If they do not, the gap is fluency, not governance. Adding more policy will not close it.
- If they do, the gap is structural — the policy and processes need redesign, not the people.
- In most cases, the gap is some mix of both. The mapping helps name which.

“The ten pillars tell you what to govern. The four competencies tell you what kind of human can govern it. Both questions have to be answered.”

CHAPTER 8

Mapping fluency to the seven lifecycle stages.

Where in an AI system's life each competency does its work.

If the previous chapter mapped fluency to the static structure of governance, this one maps it to the dynamic structure of an AI system's life. The seven stages set out in The Executive's Guide to the AI Product Lifecycle — from concept through decommissioning — are the journey every AI system in your organisation is on. Each stage requires different human capabilities. The mapping below names which.

STAGE	DOMINANT COMPETENCIES — AND WHAT GOOD LOOKS LIKE AT THIS STAGE
1. Concept & Use Case Definition	Delegation, dominant. Is AI the right tool for this problem at all? Is the problem well enough understood to delegate any of it? Strong fluency at this stage is the cheapest place to fix things — the executive sponsor with the discipline to refuse a poorly-defined use case.
2. Design & Development	Description and Diligence (creation). The team translates intent into specification — requiring strong description — and chooses data, tools, and architectures — requiring strong creation diligence. Most failures here are diligence failures dressed as technical decisions.
3. Validation & Testing	Discernment, dominant. Validation is the formalised application of organisational discernment, before deployment. The team must be able to recognise unfair, unreliable, or unsafe behaviour — not just inaccurate behaviour. Independent reviewers add discernment the build team cannot supply.
4. Deployment	Diligence (deployment). The handoff from project to operation is fundamentally a diligence question — who is now accountable, what disclosures have been made, what oversight is in place from day one.
5. Operation & Monitoring	Discernment, dominant. The Description–Discernment loop running in operational time. Without monitoring, drift is invisible; without discernment in the people reading the monitoring, drift is visible but unrecognised.
6. Maintenance & Improvement	All four, in cycle. Each maintenance change is a small lifecycle within a lifecycle — delegate the change, describe what is required, validate (discern) the result, take responsibility (diligence) for the updated system. Weak fluency here produces silent drift.
7. Decommissioning	Diligence (deployment, in reverse). The discipline to actually retire systems — protect data, revoke access, document the retirement,

STAGE	DOMINANT COMPETENCIES — AND WHAT GOOD LOOKS LIKE AT THIS STAGE
	capture lessons — is organisational diligence applied to endings. The most overlooked stage and the one most exposed to weak diligence.

Where competence shifts during the lifecycle.

Reading the mapping vertically reveals an important pattern. Delegation dominates at the front of the lifecycle. Description and creation diligence dominate in the middle. Discernment dominates in validation and operation. Deployment diligence dominates at handover and at decommissioning. Different stages require different competencies — and, in most organisations, different people. Lifecycle fluency is not about every individual being equally fluent in every competency. It is about the right competency being present at the right stage, owned by the right person.

This has practical implications for how organisations build fluency. Training every employee to use prompts well will not produce fluent validation. Training the audit team to use AI tools will not produce fluent concept-stage decisions. The pattern of fluency required is structured: certain roles need delegation fluency more than others; certain roles need discernment fluency more than others; certain roles — the executives — need all four, applied at the right level. Building organisational fluency is therefore a pattern of role-specific investments, not a uniform programme.

THREE LIFECYCLE IMPLICATIONS FOR BOARDS

- If your concept-stage delegation is weak, every later stage will be expensive — because the wrong systems will have been approved.
- If your operation-stage discernment is weak, the systems you have approved will degrade silently — the cost showing up only when it surfaces externally.
- If your decommissioning diligence is weak, your AI estate will accumulate residual systems that nobody owns and nobody has retired — and the eventual incident will involve one of them.

The board cannot oversee every stage of every AI system. It can ensure that the stages most exposed to fluency weakness are covered — by named owners, by documented processes, by evidence the board can review. Concept stage is one such point of exposure. Operation is another. Decommissioning is the third. If your board's AI oversight covers those three points well, the rest is recoverable. If it does not, the rest is at risk.

“Fluency is not a posture. It is a pattern — different in each role, different at each stage, but always present where the work is done.”

CHAPTER 9

Building organisational fluency.

A board-level roadmap. Twelve months of work. Modest in cost. Decisive in effect.

Most organisations approach AI fluency the way they approach any training need: as a procurement decision. Find a course, roll it out, measure completion, move on. That approach produces compliance with a training requirement. It does not produce fluency. Fluency is built the way any organisational capability is built — through structured exposure, role-specific investment, modelled behaviour from the top, and a feedback loop that improves the programme over time. None of that is found in a generic course.

This chapter sets out a twelve-month roadmap for building organisational fluency, designed to be achievable inside a normal management rhythm and visible to a board. It is not transformational; it is foundational. After twelve months, the organisation has the layer the previous two volumes assumed. Without that layer, the governance system the first volume described will not function as intended.

Months one to three: assessment and architecture.

The work in the first quarter is to know what you have and to design what you need. This means a fluency baseline assessment: which roles in the organisation work with AI today, what is their current capability, where are the gaps. It also means a target-state architecture: what level of fluency is required for each material role, expressed in terms of the four competencies and the three modes. The output of the quarter is a fluency map of the organisation — not a perfect one, but specific enough that the next quarter's investment is targeted.

The board's role in this quarter is to fund the assessment and to commit to the architecture. Most organisations skip this step and start with training. The training is then poorly targeted, undertaken by the wrong people, and disconnected from the governance work. A baseline assessment costs little. Skipping it costs more, eventually.

Months four to six: foundations and signal.

The second quarter is when foundational fluency is built and the organisation's posture is signalled. Foundational fluency: the universal layer — every employee who works with AI in any meaningful way receives a baseline fluency education, anchored in the four competencies. The Anthropic AI Fluency course is freely available and is a strong foundation; supplementary material may be needed for organisation-specific context.

Signalling: the executive team itself goes through fluency training, visibly, before the rest of the organisation. The board is briefed on what fluency means and how the executive team is being assessed against it. This sequencing matters. An organisation cannot rise above the fluency of its leadership; if the fluency programme starts at junior levels, it never reaches the executive table, and the governance work the previous volumes described continues to lack the human layer it needs.

Months seven to nine: role-specific deepening.

The third quarter takes the foundational layer and deepens it where the lifecycle and pillar mappings show it matters most. Concept-stage owners — typically senior executives sponsoring AI use cases — receive specific delegation fluency training. Operations leaders responsible for AI monitoring receive specific discernment fluency training. Compliance, audit, and risk professionals receive specific diligence fluency training. Procurement and IT receive specific creation-diligence training, focused on vendor and tool evaluation.

The pattern of investment is shaped by the mappings in the previous two chapters. The aim is not equal fluency across roles but appropriate fluency in each role — the right competency, at the right depth, where the lifecycle and pillar work requires it.

Months ten to twelve: integration and feedback.

The fourth quarter is when fluency stops being a programme and becomes a discipline. AI fluency is integrated into onboarding, performance reviews, and promotion criteria. The organisation begins running structured retrospectives on AI incidents and near-misses, treating each as a fluency feedback signal — what competency was missing, where, and how does the programme need to change in response. The fluency programme itself becomes part of the AI Management System the first volume described — reviewed, measured, improved, audited.

By the end of the fourth quarter, the organisation has a fluency baseline assessment, a target architecture, foundational training across the workforce, deep training in the roles that matter most, and a feedback loop that keeps the programme aligned with how AI in the organisation is actually evolving. None of this is exotic. All of it is achievable in twelve months. Most organisations have not done it.

TWELVE-MONTH FLUENCY ROADMAP — AT A GLANCE

Q1 — Assessment & architecture

- Baseline current fluency by role; design target-state architecture; secure board commitment to programme. Deliverable: fluency baseline map by role.

Q2 — Foundations & signal

- Universal foundational fluency education; executive team trained first and visibly; board briefed. Deliverable: executive team fluency assessment results and board briefing document.

Q3 — Role-specific deepening

- Targeted competency training for concept owners, operations leaders, audit/compliance, procurement — shaped by the lifecycle and pillar mappings. Deliverable: role-specific training completion records and competency assessments.

Q4 — Integration & feedback

- Fluency embedded in onboarding, performance, promotion; structured incident retrospectives; programme treated as part of the AIMS. Deliverable: fluency embedded in HR processes, first incident retrospective completed, programme review scheduled.

WHERE FLUENCY PROGRAMMES GO WRONG

- Starting with training rather than assessment — producing a programme aimed at the wrong gaps.
- Treating fluency as a junior-level requirement — leaving the executive team untrained and unable to model the discipline.
- Buying a generic course and treating completion as the goal — a metric that produces compliance, not fluency.
- Running fluency as a programme separate from governance — producing two parallel investments that don't reinforce each other.
- No feedback loop — building the programme once and not updating it as AI evolves.
- No board visibility — letting fluency become an HR initiative rather than a governance one.

THREE QUESTIONS TO ASK YOUR ORGANISATION

- Of the AI fluency we currently have in the organisation, how much was deliberately built and how much accumulated by accident?
- Is our executive team's fluency ahead of, level with, or behind that of the operational teams — and if behind, what does that mean for the credibility of our governance?
- When something goes wrong with AI in our organisation, what do we learn from it about our fluency — and where does that learning go?

CLOSING

Three commitments.

If the rest of this guide is correct, three things follow. Each is a board-level commitment.

AI fluency is not a course completion certificate. It is a posture the organisation adopts — modelled by the top, structured into the work, audited as part of the management system. Three commitments turn the argument of this guide into action.

Commitment one: govern and develop, in equal measure.

The first commitment is the central argument of this guide: governance and fluency are funded together, owned together, built together, and audited together. Governance without fluency produces theatre. Fluency without governance produces drift. The board ensures that neither investment is allowed to outrun the other — and that both are visible at executive level, not delegated to functions that will pursue one and ignore the other.

Commitment two: lead the fluency, do not delegate it.

The second commitment is that the executive team itself goes first. Fluency in an organisation rises only as high as the fluency at the top. A board that asks employees to be discerning about AI but accepts AI-generated evidence uncritically in board papers will not produce a discerning organisation. A CEO who asks employees to disclose AI involvement but conceals AI use in their own work will not produce a transparent organisation. Modelling is the most consequential investment any leadership team makes in fluency — and the cheapest.

Commitment three: treat fluency as evidence, not effort.

The third commitment is that fluency is measured, reported, and audited like any other capability the organisation depends on. Not training hours — outcomes. Not course completion — demonstrated capability. Not generic capability — the specific competencies the lifecycle and pillar mappings show this organisation needs. Fluency, like governance, is real to the extent that it can be evidenced. Treat it that way and the organisation gets the discipline it has paid for. Treat it as effort and the organisation has bought a metric.

“Two volumes told you what to govern and how. This volume tells you who has to be capable of doing it. The answer is everyone, calibrated to their role — starting with you.”

This is the third volume in the series. It does not stand alone, and was not designed to. Read alongside The Executive's Guide to AI Governance and The Executive's Guide to the AI Product Lifecycle, it completes the executive picture: what to govern, how to govern it, and who has to be capable of governing it. The next volume in the series will take a different cut at the same problem — the people problem of AI inside organisations — from the perspective of culture, communication, and change. It will arrive when it is ready.

Until then, the work is in front of every executive team that has read this far. The clock is the same clock the first volume described: South African policy moving toward implementation, ISO/IEC 42001:2023 settling as the global reference, customers and regulators beginning to ask the questions the standard is designed to answer. Building the human capability to answer those questions takes about a year. The clock is running. So is the opportunity.

As outlined earlier in this document it is good practice and an element of diligence to disclose when AI is used.

DILIGENCE STATEMENT :

In putting this guide together, I made use of Claude, Chat GPT created the images. I take full responsibility for the output of this augmentation and undertake to ensure that it is regularly updated.

Where are you now? A four-competency self-assessment.

For each of the four AI Fluency competencies, mark your organisation's current position honestly — both for the executive team and for the broader workforce. This is a diagnostic, not a test. Where you mark Not in place is where to focus first.

Competency	Exec team: In place	Exec team: Partial	Workforce: In place	Workforce: Partial
Delegation — which categories of work AI is permitted to do, and which are off the table	☐	☐	☐	☐
Description — clear policies, communications, and briefs that shape how AI is used across the organisation	☐	☐	☐	☐
Discernment — structured review of AI outputs, monitoring of AI performance, culture of questioning AI evidence	☐	☐	☐	☐
Diligence — named accountability, disclosure mechanisms, and incident response for AI-mediated outcomes	☐	☐	☐	☐

If the executive team column is weaker than the workforce column, the fluency programme needs to start at the top — not roll out bottom-up. If Discernment and Diligence are the weakest competencies across both columns, the governance system is at risk regardless of how good the policies are. Use the Chapter 7 mapping to identify which governance pillars are most exposed by the gaps you have found.

For a facilitated version of this assessment with role-level benchmarking, take the AI Fluency Assessment at <https://theaifluencyco.com> or book a 30-minute conversation at <https://calendly.com/zoe francois123/30min>.

ABOUT

The AI Fluency Company.

The AI Fluency Company helps executive teams adopt AI responsibly — building the language, the judgement, and the governance to make AI work in their business. We work with mid-market organisations, family businesses, and compliance-minded leaders who want to do AI properly the first time.

How we work with executive teams.

ENGAGEMENT	WHAT IT IS
AI Fluency Briefing	For Leadership teams who know AI matters, but they do not yet share language, boundaries, decision criteria, or a practical view of what AI should and should not do in the organisation. This is a half-day workshop for leadership teams. Builds shared understanding of working with AI effectively, efficiently, ethically, and safely.
AI Governance Readiness Diagnostic	Many organisations have AI activity already underway in their businesses. Tools, vendors, pilots, staff experimentation, and automated decisions, are in place but ownership, controls, risk, and value are not being accounted for.
AI Strategy Workshop	The organisation has interest, ideas, and perhaps pilots, but lacks a disciplined way to evaluate and prioritise AI opportunities, size value, assign ownership, and decide what should move forward. This two day workshop provides methodology to prioritises the use cases, map ownership, build roadmaps.
Hyperautomation Opportunity Sprint	Teams are using AI to speed up fragments of work, but the larger workflow remains slow, manual, fragmented, or badly designed. The value is trapped in the process, not the tool. Structured engagement to apply the principles in this playbook.
AI Governance Implementation Retainer	If after the diagnostic or workshop, the organisation requires sustained help translating recommendations into policy, controls, ownership, reporting, and decision habits a retainer option is available.

About the author.

Zoë François is the founder of The AI Fluency Company. She trained as a chemical engineer and built one of South Africa's earliest applied industrial AI systems in the 1990s — an expert system on the Gensym G2 platform that modelled Anglo American Platinum's refinery operations. She then spent two further decades in commercial leadership, ultimately as **General Manager of Sales at Sappi** and **Senior Sales Executive at Consol Glass**. She holds the Oxford Saïd Business School AI Programme certificate.

Engineering rigour. Commercial judgement. The combination most AI engagements are missing.

TAKE THE NEXT STEP

Take the AI Fluency Assessment for a maturity benchmark: <https://theaifluencyco.com>

Book a 30-minute conversation to walk through your AI fluency starting point. No charge for the first call.

Read the companion volumes — The Executive's Guide to AI Governance and The Executive's Guide to the AI Product Lifecycle — for the full ten-pillar framework and the seven-stage lifecycle.

Take the underlying course on which Chapter 2 is based: AI Fluency: Framework and Foundations, free at anthropic.com/ai-fluency. Authored by Rick Dakan and Joseph Feller in partnership with Anthropic, published 2025 under CC BY-NC-SA 4.0.

First edition · 2026 · Cape Town

© The AI Fluency Company. All rights reserved.