THE AI FLUENCY CO

EVALUATE. ELEVATE. EMPOWER.

TRANSFORMATION SERIES

The Hyperautomation Playbook

Don’t automate the old process.

Reconceptualise it.

From automation to transformation.

By Zoë François

Founder, The AI Fluency Company

First edition · 2026

CONTENTS

What's in this playbook.

FOREWORD

Why this playbook exists.

AI without process redesign is a slightly faster broken workflow. This is a guide for the leaders who refuse to settle for that.

This is the first volume in The Transformation Series — a parallel set of guides from The AI Fluency Company aimed at the people whose job is to deliver business value from AI. If our Executive Series is for the people who govern AI, this series is for the people who use it to transform what their organisation actually does.

The motivation for this playbook is a pattern I have watched accumulate across dozens of conversations with operations leaders, transformation directors, and chief operating officers. The pattern goes like this. The organisation invested in AI, the pilots ran, vendors were paid dashboards were built. And the business outcomes — cycle time, cost, customer experience, throughput — are unchanged or marginally better. Something the executive sponsor cannot quite name has gone wrong, and the conversation usually ends with some version of ‘we need a better AI strategy’ or ‘we need different tools’ or ‘we need better people’.

None of those is the answer. The answer is harder and more interesting. The reason most AI programmes underdeliver is that the organisation automated its existing processes rather than reconceptualising them. They took workflows designed for a world without AI — designed for paper, for handoffs between specialists, for sequential approval chains, for human pattern-recognition — and they layered AI on top. The result is what McKinsey’s researchers have consistently described as a slightly faster broken workflow. The bottleneck did not disappear. It moved.

> *“If you take a broken, manual, approval-heavy workflow and add AI on top of it, you get a slightly faster broken workflow. The bottleneck doesn’t disappear — it moves.”*

This is not a hypothetical concern. McKinsey’s 2025 State of AI survey, conducted across nearly two thousand companies in 105 countries, tested a range of organisational attributes and found that workflow redesign had the single strongest correlation with EBIT impact from AI. The high-performing organisations — roughly six per cent of the sample — were nearly three times more likely to have fundamentally redesigned their workflows around AI rather than layered AI onto existing processes. The same pattern shows up in BCG’s research. The same pattern shows up in every serious analysis of why digital transformation programmes fail — and seventy per cent of them do, by Boston Consulting Group’s reckoning, with most of the rest underperforming their original ambitions.

This playbook is the operating manual for the other path. It is for the leader who has read the same news articles, watched the same vendor pitches, heard the same productivity claims — and decided that productivity is not the same thing as transformation. The leader who knows that automating a finance reconciliation that should not exist is worse than not automating it at all. The leader who is willing to ask the harder question: not ‘how do we make this process faster with AI?’ but ‘what would this process look like if we designed it knowing AI was available?’

That question is the difference between automation and hyperautomation, and it is the spine of this playbook. The first three chapters make the argument: that the value from AI is not in the tools, it is in the redesign; that hyperautomation is the discipline of orchestrated reconceptualisation; that organisations which underinvest in process redesign systematically underperform. The next four chapters give you a way to do the work: a seven-step playbook for moving from current state to transformed state, with the same callout-box discipline as our Executive Series — what each step requires, what good looks like, where organisations go wrong, three questions to ask. The closing chapters connect this work to the governance and fluency disciplines covered in our other series, because transformation done without governance creates risk, and transformation done without fluency creates failure.

A note on what this guide is not. It is not a vendor evaluation. It is not a tool selection guide. It is not a deep-dive into RPA, process mining, or any single technology category — those exist in abundance elsewhere. It is a leadership document, written for the person who will sponsor a transformation programme and be held accountable for whether it delivers value. The technical detail is the easy part of hyperautomation. The hard part is the discipline to redesign work, sustain the redesign through delivery, and protect it from the gravitational pull of ‘let’s just automate what we have’. That discipline is what this playbook is about.

The South African context, as in our other guides, matters. The local economy needs transformation programmes that actually deliver — productivity gains, cost reductions, capacity created for higher-value work, sustainable competitiveness against international peers. The technologies are available. The frameworks are settled. What is missing, in most organisations I work with, is the leadership posture this playbook describes: the willingness to do the hard reconceptualisation work before reaching for the tools.

Zoë François

Cape Town, 2026

CHAPTER 1

The reconceptualisation imperative.

The single most consequential decision in any AI programme is made before the technology is chosen.

Every organisation I work with eventually arrives at the same fork in the road, often without realising they have arrived at it. The fork looks like this. The team has identified a process worth improving — customer onboarding, claims handling, invoice processing, demand forecasting, employee provisioning, supplier reconciliation, whichever it is. AI tools exist that could help. The vendor pitches confidently a pilot project is scoped. At this moment, before any code is written, the organisation makes a choice that will determine whether the entire programme succeeds or fails.

The choice is between two questions. The first question is: 'how do we use AI to make this process faster?' The second question is: 'what would this process look like if we designed it today, knowing AI is available?' The options sound similar however they produce dramatically different outcomes. Most organisations ask the first question without realising the second exists, and they spend the next eighteen months delivering a programme that automates the inefficiencies of a workflow that should not have survived contact with the technology in the first place.

The trap of automating the old process.

The reason organisations default to the first question is structural, not foolish. The current process is documented. The current process owners can describe it. The current systems support it. The current people know how to operate it. Automating the current process is the path of least resistance — it is the version of the programme that requires the fewest decisions, the fewest stakeholders, the smallest behaviour change. It is also the version with the smallest possible upside, because the upside is constrained by the design of the underlying process, which was itself a compromise built around the possibilities of an earlier era.

Consider a typical mid-market customer onboarding process. The customer fills in a form. The form goes to a queue. A staff member reviews it for completeness. Information is verified across three or four systems. A risk assessment is generated. An approval moves through one or two layers of management. Compliance does a separate check. The customer is notified. End-to-end, this takes between three and fifteen days, depending on the organisation. AI can be applied to almost every step — to validate the form, to extract data, to score risk, to draft the approval, to write the customer notification. Doing all of that produces a process that takes between one and five days, with fewer human touchpoints.

That is automation it is not transformation. The transformed version of the same process does not look like the original process running faster. It looks like a process that did not exist before. The customer is onboarded in the moment they apply, with AI handling identification, validation, risk scoring, and decisioning in real time. Edge cases are routed to humans, who handle them with full context generated by the AI. The process is not a sequence of handoffs anymore — it is an event. The end-to-end time is minutes, not days. The staff who used to operate the queue now manage exceptions and improve the model. The cost per onboarding is a fraction of what it was, but more importantly, the customer experience is different in kind, not in degree.

> *“Automation makes the existing process faster. Transformation makes the existing process unnecessary.”*

The McKinsey finding.

This pattern is not anecdotal. McKinsey’s 2025 State of AI survey, conducted across nearly two thousand organisations in 105 countries, tested a range of organisational attributes against the question of which ones correlated with measurable EBIT impact from AI. Workflow redesign had the strongest correlation of any attribute tested — stronger than tool selection, stronger than data quality, stronger than executive sponsorship, stronger than budget allocated. The high-performing organisations were nearly three times more likely to have fundamentally redesigned their workflows around AI than to have layered AI onto existing processes.

What that finding actually says is not that workflow redesign helps, it is that workflow redesign is the single strongest predictor of whether an AI programme produces financial returns. Organisations that redesign their workflows are roughly three times more likely to be in the small group that gets meaningful financial returns from AI. Organisations that do not redesign their workflows are, on average, in the larger group that does not. The difference between value and no-value is, more than any other variable, the difference between automation and reconceptualisation.

This is the single most important strategic finding in AI today. It also goes substantially against the grain of how most organisations are currently approaching AI. The default mode is tool-first: pick the AI capability, find a process to apply it to, measure the productivity gain. The mode that produces value is process-first: pick the outcome that matters, redesign the process to deliver that outcome with AI as a fundamental component, then choose the tools that fit. The order of operations matters. Reversing it does not produce slightly worse results. It produces, on average, no results.

What reconceptualisation means in practice.

Reconceptualisation is not a synonym for ambition. Many organisations describe their AI programmes as ambitious and still produce slightly-faster-broken-workflows. The defining feature of reconceptualisation is the willingness to ask, of every step in the existing process, the question ‘would this step exist if we were designing this process today?’ Most steps in most processes exist for reasons that no longer apply. They were inserted because a system could not communicate with another system, or because a regulator required a particular form of evidence in a particular sequence, or because a manager wanted oversight at a particular point, or simply because a different person was responsible for the next stage. Each of those reasons may or may not still apply. Most do not.

Reconceptualisation is the discipline of asking that question seriously, of every step, and being willing to remove the steps that do not justify themselves. It is also the discipline of asking the inverse question — 'what new things become possible because AI is in the process?' — and being willing to add capabilities that did not exist in the original design. A reconceptualised process typically has fewer steps, different roles, different metrics, different cycle times, and a different value proposition to the customer or end user. It is a different process, not a faster version of the same process.

WHAT RECONCEPTUALISATION REQUIRES

- A willingness to question every step in the current process — not to defend it.
- A clear view of what AI now makes possible — not constrained by what was previously feasible.
- An executive sponsor with the authority to remove steps, change roles, and rewrite metrics.
- A target-state design produced before tool selection, not derived from it.
- A change-management plan proportionate to the magnitude of the redesign — not bolted on at the end.
- Discipline to resist the gravitational pull of 'let's just automate what we have for now'.

WHERE ORGANISATIONS DEFAULT TO AUTOMATION INSTEAD OF TRANSFORMATION

- Tool-first thinking — starting with the AI capability and looking for a process to apply it to.
- Letting the vendor scope the redesign — producing a transformation that is conveniently the size of the vendor's product.
- Process redesign delegated to the team that owns the current process — rarely producing fundamental change.
- Treating change management as an after-the-fact communications exercise rather than as the substrate of the redesign.
- Stopping at the first version that 'works' — banking the productivity gain and never returning to the deeper redesign.
- Confusing pilot success with programme success — a pilot proves the technology, not the transformation.

THREE QUESTIONS TO ASK YOUR ORGANISATION

- Of the AI initiatives currently funded in our organisation, how many are automating an existing process — and how many are redesigning it?

- Could we name a process in our business that we have fundamentally reconceptualised in the last three years — and what changed about it, beyond speed?
- If our most successful AI programme delivered the result it is currently designed to deliver, what fraction of the available value would we have captured?

If those questions are uncomfortable, you are in the same position as most organisations — not because you have failed, but because the path of least resistance has done what paths of least resistance always do. The next two chapters look more closely at what hyperautomation actually is, and at the systemic reasons most AI programmes underdeliver. From Chapter 4, the playbook itself begins.

CHAPTER 2

From automation to hyperautomation.

Four words that look similar. Four practices that produce very different outcomes.

The word 'automation' has been used to mean too many different things, by too many different vendors, for too many different audiences. The result is a market in which executives are not sure what they are buying, technical teams are not sure what they are delivering, and boards are not sure what they have approved. The first thing this chapter does is restore some discipline to the vocabulary.

There are, in practice, four distinct stages in the evolution of automation, and they differ from each other not in degree but in kind. The fourth — hyperautomation — is the only one that consistently produces the order-of-magnitude gains executives associate with AI transformation. The first three each have their place, but mistaking one for another is the source of most disappointment in AI programmes.

The four stages, named precisely.

STAGE	WHAT IT DOES	WHAT IT CANNOT DO
1. Task automation	Automates a single, well-defined, repetitive task. Rules-based. Stable inputs, stable outputs.	Cannot handle exceptions. Cannot adapt. Cannot orchestrate beyond the task. Doesn't 'think'.
2. Intelligent automation	Adds AI/ML to the task — unstructured data, classification, prediction, judgement on individual items.	Still task-bound. Doesn't redesign the workflow it sits inside. Smarter component, same process.
3. Process automation	Orchestrates multiple tasks across systems and people — end-to-end workflow execution.	Automates the existing workflow. Doesn't question whether the workflow itself should exist.
4. Hyperautomation	Reconceptualises the work itself — redesigns processes around what is now possible, orchestrates AI, automation, data, and people as one system, continuously improves.	Requires real organisational discipline — leadership, redesign authority, change management and governance. Not a tool you buy.

The taxonomy matters because the four stages produce dramatically different returns, require dramatically different investments, and depend on dramatically different organisational capabilities.

An organisation that buys hyperautomation tooling and applies it as task automation produces the returns of task automation. An organisation that approaches task automation with the discipline of hyperautomation produces returns that are out of proportion to the technology involved. The defining variable is not the tool. It is the level of redesign the organisation is willing to do.

The Gartner definition, and what it actually means.

Gartner, which coined the term, defines hyperautomation as ‘a business-driven, disciplined approach that organisations use to rapidly identify, vet and automate as many business and IT processes as possible’. The definition involves orchestrating multiple technologies — AI, machine learning, RPA, business process management, low-code platforms, integration tools, and analytics — as a coordinated system rather than as separate point solutions.

Two phrases in the Gartner definition deserve more attention than they usually get. The first is ‘business-driven’. Hyperautomation, done correctly, is not a technology programme that the business is asked to support — it is a business programme that uses technology. The distinction sounds rhetorical. It is not. A technology-led programme produces technology outcomes — systems deployed, integrations built, models trained. A business-led programme produces business outcomes — cycle time reduced, costs lowered, customer experience changed. Most failed hyperautomation programmes were technology-led programmes that mislabelled themselves.

The second phrase is ‘disciplined approach’. Hyperautomation is not a procurement event. It is an ongoing practice, with its own governance, its own metrics, its own talent requirements, its own funding cycle. Organisations that treat it as a one-off transformation project deliver one-off transformation outcomes — a brief surge of value followed by gradual decay as the redesigned processes drift back toward their original shape. Organisations that treat hyperautomation as a permanent organisational capability deliver compounding value over years.

> *“Hyperautomation is not a tool you buy. It is a discipline you build.”*

Why the distinction is a leadership issue.

It would be tempting to leave the distinction between automation and hyperautomation as a vocabulary point and move on. That would be a mistake, because the confusion between them is the most common source of executive misalignment in transformation programmes. The CFO believes the organisation is investing in cost reduction. The COO believes it is investing in capacity creation. The CIO believes it is investing in technology modernisation. The CEO believes it is investing in transformation. Each is technically correct about a different one of the four stages. None of them is talking about the same programme.

The first responsibility of any transformation leader is to establish, with the rest of the executive team, which of the four stages the organisation is actually pursuing — and to be honest about it. There is no shame in pursuing task automation; it is the right answer for many problems. There is also no shame in pursuing hyperautomation; it is the right answer for the largest opportunities. The shame is in pursuing one and pretending it is the other, which is what most organisations are doing today.

DIAGNOSTIC: WHICH STAGE IS YOUR PROGRAMME ACTUALLY AT?

- If most of your AI work is rules-based automation of single tasks, you are at stage 1. The returns are real but bounded — typically 10–20 per cent productivity gain on the automated task.
- If your AI work adds intelligence to individual tasks (classification, prediction, document processing), you are at stage 2. Returns are larger — typically 20–40 per cent on the automated activities — but still task-bound.
- If your AI work orchestrates end-to-end workflows across systems, you are at stage 3. Returns can reach 30–50 per cent on the automated processes, but you have not yet questioned whether the processes themselves should exist.
- If your AI work is redesigning the work itself — not just executing it faster — you are at stage 4. This is where the genuinely transformational outcomes live. Gartner has projected that organisations combining hyperautomation technologies with redesigned processes can lower operational costs by 30 per cent across the board — not just on the automated activities, but organisation-wide.

The diagnostic above is uncomfortable for most organisations because it reveals that what they have been calling transformation is, in fact, a more mature version of automation. That recognition is the entry point to the work this playbook describes. The next chapter looks at the structural reasons most AI programmes stall at stage three, and what would have to change for them to reach stage four.


THE AI FLUENCY CO
EVALUATE. ELEVATE. EMPOWER.
THE SEVEN FAILURE MODES
THAT HINDER HYPERAUTOMATION SUCCESS
Common patterns that derail initiatives, waste investment and prevent real business value.
Avoiding these failure modes is the fastest path to sustainable transformation.
01
TOOL-FIRST THINKING
Starting with capability rather than business outcomes and value.
RESULT:
Solutions that don't solve the right problems.
02
AUTOMATING THE EXISTING PROCESS
Automating the current process — not reconceptualising it.
RESULT:
Manual inefficiencies are simply automated at scale.
03
PILOT PURGATORY
Brilliant pilots that don't scale beyond a small sandbox.
RESULT:
Value is never realised beyond isolated use cases.
04
CHANGE MANAGEMENT AS AFTERTHOUGHT
Communication rather than adoption, culture and capability.
RESULT:
Low adoption, resistance and failure to embed change.
05
DATA UNDERESTIMATED
Quality, access and governance not addressed early.
RESULT:
Poor data leads to poor decisions and loss of trust.
06
GOVERNANCE BOLTED ON
Governance, risk and compliance treated as an add-on.
RESULT:
Risk exposure, compliance issues and rework.
07
NO CONTINUOUS IMPROVEMENT
Transformation treated as a project, not an ongoing capability.
RESULT:
Value plateaus, opportunities missed, and momentum lost.
THREE QUESTIONS TO ASK YOUR ORGANISATION
?
01
Of the seven failure modes above, how many are present in our current programme — and which is the one we are most reluctant to admit?
?
02
When did we last cancel an AI initiative because it was reproducing one or more of these patterns — or do we tend to push through and rationalise the result afterwards?
?
03
Who in the organisation is empowered to call out these patterns when they appear — and is that authority real, or theoretical?
RECOGNISE. ADDRESS. TRANSFORM.
Success comes not from avoiding change, but from avoiding the patterns that cause failure.

CHAPTER 3

Why most AI programmes underdeliver.

Seven failure modes. Each preventable. Each common.

The headline failure rate for digital transformation programmes is seventy per cent, and has been for a decade, across every major management research firm that tracks it. The failure rate for AI-specific transformations is, on current evidence, no better. Bain's 2024 analysis put the figure at eighty-eight per cent of business transformations failing to achieve their original ambitions. McKinsey has consistently found that only a small minority — roughly six per cent in their 2025 State of AI survey — of organisations report meaningful financial returns from their AI investments. The numbers are sobering, but they are also informative, because the failure modes are well-documented and consistent across studies.

This chapter sets out the seven failure modes I see most often in mid-market and enterprise transformation programmes. Each is preventable. Each is common. Most failed programmes exhibit at least three of the seven simultaneously. If you are sponsoring or leading a transformation programme, the test of this chapter is whether you can read it without recognising at least two of the patterns in your own organisation.

Failure mode 1: tool-first thinking.

The most common failure pattern starts with a vendor demonstration or an industry conference, not with a business problem. Someone in the organisation becomes excited about a capability — generative AI, autonomous agents, intelligent document processing — and starts looking for places to apply it. The capability is procured. Use cases are scoped to fit the capability. The pilot succeeds because pilots are designed to succeed. And then the programme stalls when it tries to scale, because the organisational change required to make the capability deliver value at scale was never planned for. The capability was the answer; nobody had asked the question.

Failure mode 2: automating the existing process.

Covered in detail in Chapter 1, this is the largest single source of underperformance in AI programmes. The team takes the existing process, identifies the steps where AI can be applied, and produces a faster version of the same process. The headline metrics improve modestly. The business outcomes — cost, capacity, customer experience — do not change in kind. The team reports success because the metrics did improve. The board approves the next programme on the same basis. The cycle continues until someone notices that twelve months and several million rand of investment have produced an organisation that is, in every meaningful respect, the same organisation it was before.

Failure mode 3: pilot purgatory.

The third failure mode is the pilot that never scales. Every transformation conference includes case studies of brilliant pilots. Far fewer include case studies of brilliant pilots that became enterprise capabilities. The reason is that piloting and scaling require different organisational disciplines, and most organisations are good at one and bad at the other. Piloting rewards experimentation, narrow scope, dedicated teams, and tolerance for failure. Scaling rewards standardisation, broad coordination, embedded teams, and tolerance for friction. The skills that produced the pilot success are often the wrong skills for the scaling work, and the handover is rarely planned for.

Failure mode 4: change management as afterthought.

Most transformation programmes treat change management as a workstream that runs in parallel with the technology delivery, often led by a different team, with separate budget and separate reporting. This produces a predictable failure: by the time the change management programme reaches the people whose work is being redesigned, the redesign has already been decided, and the role of change management is reduced to communication. Communication is not change management. Real change management starts with the people whose work is being redesigned being central to the redesign — not informed of it after the fact. Programmes that get this wrong produce technically successful deployments that fail in adoption, which is to say, fail completely.

Failure mode 5: data underestimated.

AI runs on data. The quality, accessibility, and governance of an organisation's data is the rate-limiting factor on what AI can do, and the most-underestimated dimension of every transformation programme I have seen. Organisations consistently overestimate the quality of their data and underestimate the work required to make it AI-ready. The result is programmes that arrive at the integration phase and discover that the data they need either does not exist, is not accessible, is not clean enough to use, or is governed in ways that prohibit the use case. This discovery typically happens nine months into a twelve-month programme. The programme then either delivers a much smaller capability than planned, or the programme's scope creeps, or is quietly closed.

Failure mode 6: governance bolted on.

Transformation programmes that ignore governance until late produce systems that cannot be audited, cannot be defended to regulators, cannot be explained to customers, and cannot survive their first material incident. Governance retro-fitted onto an operational AI system is significantly more expensive than governance designed in from the start, and significantly less effective. This failure mode is increasingly visible because the regulatory environment is moving — ISO/IEC

42001:2023 is being adopted globally, the EU AI Act is in force, and South African policy is targeted for implementation by 2027 to 2028. The South African context offers a pointed illustration: the Draft National AI Policy, gazetted on 10 April 2026 (Government Gazette Notice 3880 of 2026), was withdrawn just sixteen days later after journalists discovered that at least six of its 67 academic citations had been fabricated by AI that was used without adequate human oversight. A government unable to govern its own AI-assisted documents is a government that understands, viscerally, why governance cannot be bolted on. Organisations whose transformations were not designed with governance built in are about to discover the cost of that omission.

Failure mode 7: no continuous improvement.

The seventh failure mode is the one that turns successful transformations into temporary ones. The programme launches. The metrics improve and the team disbands. Six months later, the redesigned process has drifted back toward its original shape — because the organisation around it did not change, because the metrics that drive behaviour did not change, because nobody owns the continued evolution of the redesigned work. Hyperautomation is, by definition, a continuous discipline. Programmes that treat it as a one-off project deliver one-off results.

THE SEVEN FAILURE MODES — A SELF-ASSESSMENT

- Tool-first thinking — starting with capability rather than business outcome.
- Automating the existing process — not reconceptualising it.
- Pilot purgatory — brilliant pilots that don’t scale.
- Change management as afterthought — communication rather than co-design.
- Data underestimated — quality, access, and governance not addressed up front.
- Governance bolted on — not designed in.
- No continuous improvement — transformation as a project, not a discipline.

THREE QUESTIONS TO ASK YOUR ORGANISATION

- Of the seven failure modes above, how many are present in our current transformation programme — and which is the one we are most reluctant to admit?
- When did we last cancel an AI initiative because it was reproducing one of these patterns — or do we tend to push through and rationalise the result afterwards?
- Who in the organisation is empowered to call out these patterns when they appear — and is that authority real, or theoretical?

The seven failure modes are not independent. They cluster. Organisations that exhibit tool-first thinking almost always also exhibit pilot purgatory and change management as afterthought — because all three are downstream of the same root cause, which is treating transformation as a technology problem. Organisations that automate the existing process almost always also fail on continuous improvement — because both are downstream of treating transformation as a finite project. Recognising the clusters helps. Fixing the root causes is what the rest of this playbook is for.

> *"The reason most transformations fail is not a single bad decision. It is a sequence of small, reasonable decisions that compound into a programme nobody intended."*

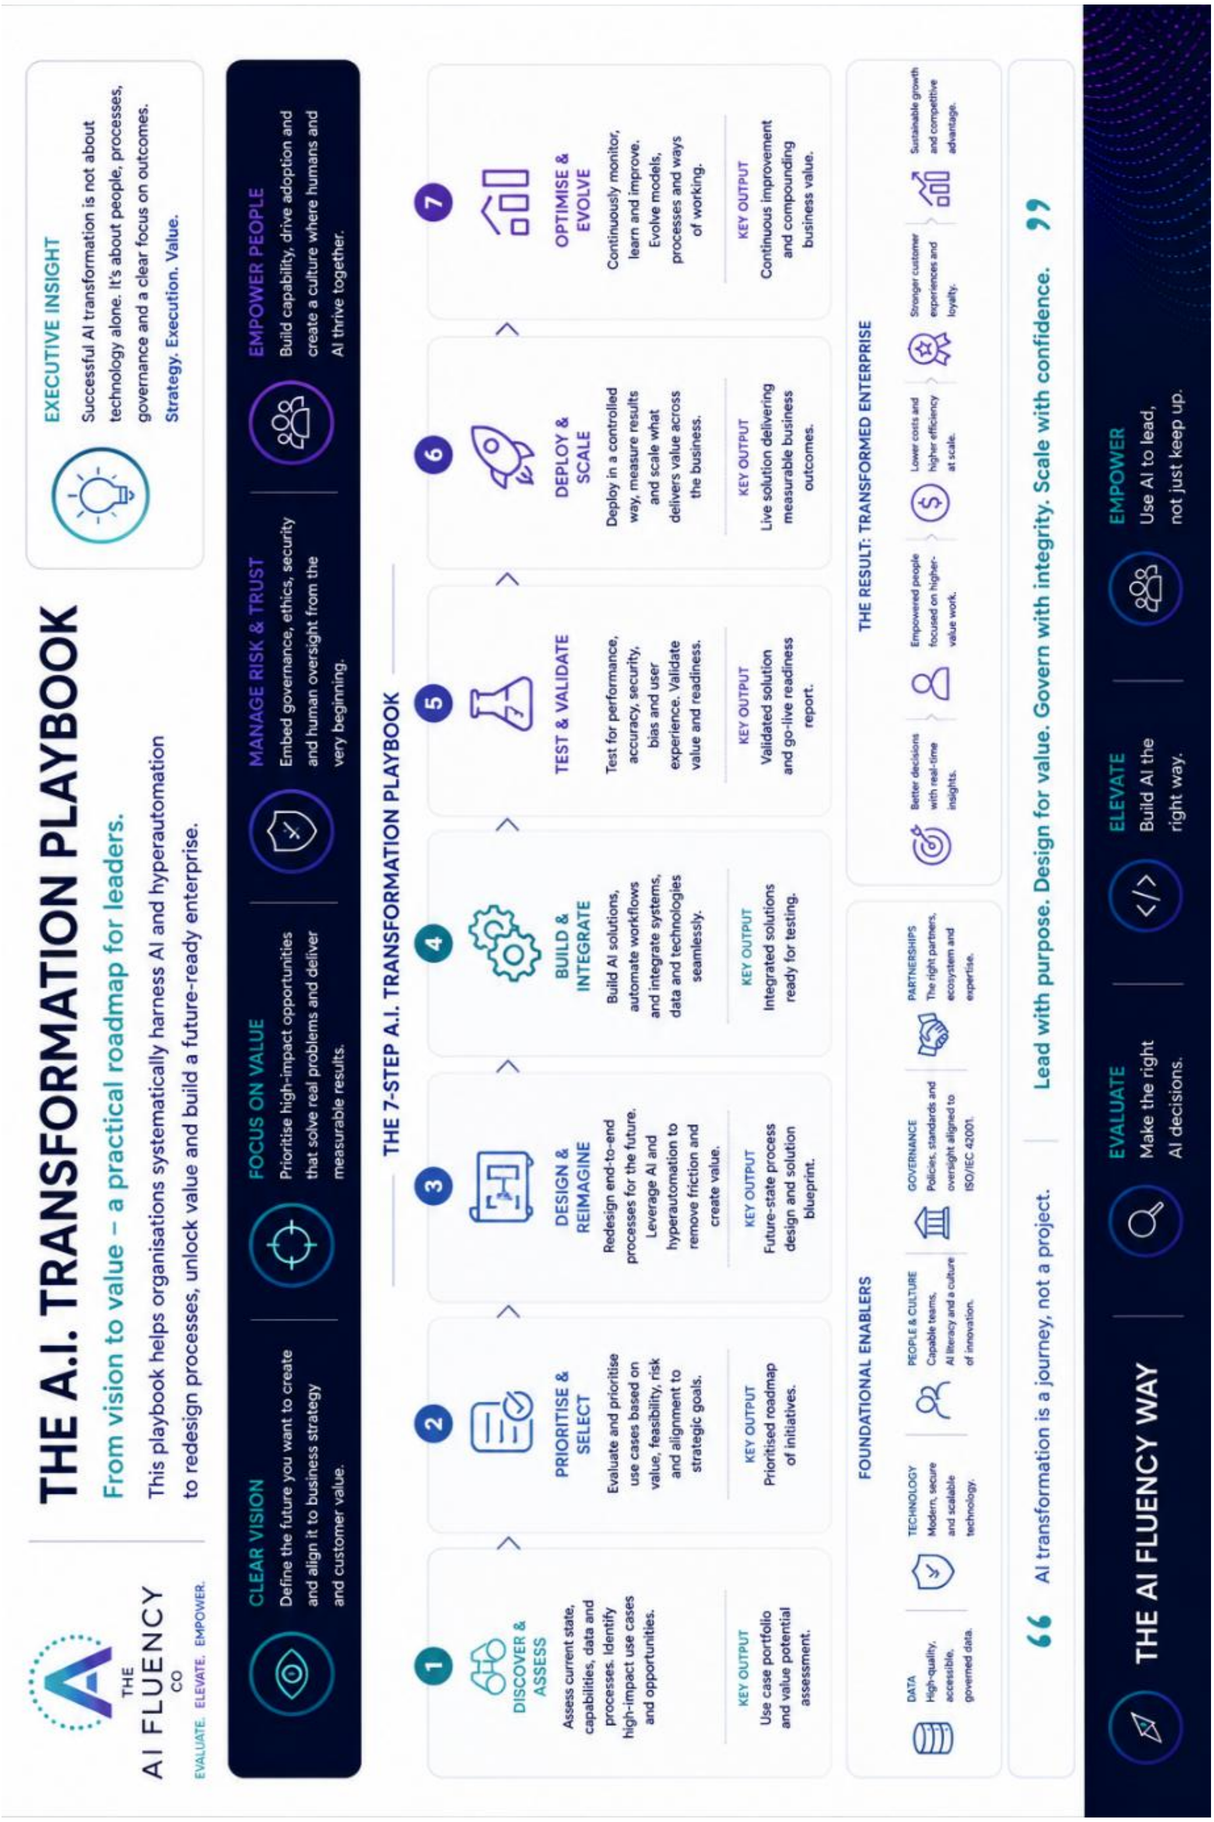
THE AI FLUENCY CO
EVALUATE. ELEVATE. EMPOWER.
THE A.I. TRANSFORMATION PLAYBOOK
From vision to value – a practical roadmap for leaders.
This playbook helps organisations systematically harness AI and hyperautomation to redesign processes, unlock value and build a future-ready enterprise.
EXECUTIVE INSIGHT
Successful AI transformation is not about technology alone. It's about people, processes, governance and a clear focus on outcomes.
Strategy. Execution. Value.
CLEAR VISION
Define the future you want to create and align it to business strategy and customer value.
FOCUS ON VALUE
Prioritise high-impact opportunities that solve real problems and deliver measurable results.
MANAGE RISK & TRUST
Embed governance, ethics, security and human oversight from the very beginning.
EMPOWER PEOPLE
Build capability, drive adoption and create a culture where humans and AI thrive together.
THE 7-STEP A.I. TRANSFORMATION PLAYBOOK
1
DISCOVER & ASSESS
Assess current state, capabilities, data and processes. Identify high-impact use cases and opportunities.
KEY OUTPUT
Use case portfolio and value potential assessment.
2
PRIORITISE & SELECT
Evaluate and prioritise use cases based on value, feasibility, risk and alignment to strategic goals.
KEY OUTPUT
Prioritised roadmap of initiatives.
3
DESIGN & REIMAGINE
Redesign end-to-end processes for the future. Leverage AI and hyperautomation to remove friction and create value.
KEY OUTPUT
Future-state process design and solution blueprint.
4
BUILD & INTEGRATE
Build AI solutions, automate workflows and integrate systems, data and technologies seamlessly.
KEY OUTPUT
Integrated solutions ready for testing.
5
TEST & VALIDATE
Test for performance, accuracy, security, bias and user experience. Validate value and readiness.
KEY OUTPUT
Validated solution and go-live readiness report.
6
DEPLOY & SCALE
Deploy in a controlled way, measure results and scale what delivers value across the business.
KEY OUTPUT
Live solution delivering measurable business outcomes.
7
OPTIMISE & EVOLVE
Continuously monitor, learn and improve. Evolve models, processes and ways of working.
KEY OUTPUT
Continuous improvement and compounding business value.
FOUNDATIONAL ENABLERS
DATA
High-quality, accessible, governed data.
TECHNOLOGY
Modern, secure and scalable technology.
PEOPLE & CULTURE
Capable teams, AI literacy and a culture of innovation.
GOVERNANCE
Policies, standards and oversight aligned to ISO/IEC 42001.
PARTNERSHIPS
The right partners, ecosystem and expertise.
THE RESULT: TRANSFORMED ENTERPRISE
Better decisions with real-time insights.
Empowered people focused on higher-value work.
Lower costs and higher efficiency at scale.
Stronger customer experiences and loyalty.
Sustainable growth and competitive advantage.
AI transformation is a journey, not a project.
Lead with purpose. Design for value. Govern with integrity. Scale with confidence.
THE AI FLUENCY WAY
EVALUATE
Make the right AI decisions.
ELEVATE
Build AI the right way.
EMPOWER
Use AI to lead, not just keep up.

CHAPTER 4

The seven-step playbook.

From discovery to compounding value — a sequence that protects against every failure mode in the previous chapter.

The seven steps below are the operational spine of a hyperautomation programme. Each step has its own purpose, its own deliverables, its own owner, and its own decision gate. The sequence matters. Steps cannot be safely skipped or reordered, and the temptation to skip the early ones — because they feel like preparation rather than action — is the single most common cause of late-stage programme failure. The four chapters that preceded this one made the case for why the work is necessary. This chapter is how the work is done.

Each step is presented in the same structure used by our Executive Series guides: what the step requires, what good looks like, where organisations go wrong, and three questions to ask. Read the seven steps as a single integrated sequence; they are not independent of each other.

Step 1: Discover & Assess.

What do we have, and where is the value?

Discovery is where the programme decides what it is for. Most transformation programmes skip this step or do it superficially — a few workshops, a list of pain points, a vendor-supplied assessment of 'AI opportunities'. That superficial discovery produces superficial programmes. Real discovery is a structured exercise in understanding the current state of the organisation's processes, capabilities, data, and pain points — and then mapping that current state to the places where AI-enabled redesign would produce the largest measurable value.

Discovery is also the only stage at which the organisation can honestly assess what it does not know. Most organisations do not know how their core processes actually run — they know how the processes are documented to run, which is different. They do not know how much time their people spend on rework versus original work. They do not know which of their data assets are actually accessible, current, and usable. The discovery stage surfaces these gaps before they become programme-stopping discoveries late in delivery.

WHAT DISCOVERY REQUIRES

- A current-state map of material business processes — not the documented version, the actual one.
- An honest capability assessment — data quality, technology stack, talent, governance maturity.
- Process mining or equivalent quantitative analysis where possible — not just interview-based mapping.
- A long list of candidate use cases drawn from the actual pain points, not from vendor catalogues.
- Initial sizing of value potential for each candidate — measured in business outcomes, not technology metrics.

WHAT GOOD LOOKS LIKE

- The discovery output is a use-case portfolio with quantified value potential — cost, cycle time, customer experience — not a wish list.
- The portfolio includes processes the organisation is bad at, not just processes the organisation is interested in.
- Capability gaps are documented honestly, including the ones that will be inconvenient to address.
- The assessment is owned by an executive sponsor with the authority to act on what it reveals.
- The output is a living document — revisited as the programme proceeds, not filed away after the kickoff.

WHERE ORGANISATIONS GO WRONG

- Discovery delegated to vendors — producing a use-case list shaped by what the vendor sells.
- Pain points captured through interviews only — missing the processes nobody wants to talk about.
- Capability gaps glossed over — producing a programme that runs into them as surprises.
- Skipping discovery because the team ‘already knows’ the processes — producing a programme based on assumptions that turn out to be wrong.
- Treating discovery as a one-off event — not as the entry point to a continuous practice.

THREE QUESTIONS TO ASK

- Could we hand someone the current-state map of our top three processes — and would it match how the work actually happens this week?
- How honest are we about which of our data assets are AI-ready — and which are not?
- If we ranked our candidate AI use cases by quantified value, how many of them are there because they’re valuable — and how many because they’re convenient?

Step 2: Prioritise & Select.

Which of the things we could do, should we do first?

A discovery process worth its name produces more candidate use cases than the organisation can possibly execute in a year. Prioritisation is the discipline of choosing among them. It is also the stage at which most programmes go wrong in a particular and predictable way: by selecting the wrong portfolio shape. The temptation is to pick a balanced mix of ‘quick wins’ and ‘transformational bets’. The reality is that the quick wins consume disproportionate organisational attention and produce disproportionately small returns, while the transformational bets are quietly underfunded and never reach the scale where they would have paid back.

Strong prioritisation is honest about the trade-offs and explicit about what is being deprioritised. Weak prioritisation says yes to everything that scored above a threshold and lets the programme team work out the sequencing. The first produces programmes that move forward; the second produces programmes that fragment.

WHAT PRIORITISATION REQUIRES

- A scoring model that includes value potential, feasibility, risk, and strategic alignment — weighted appropriately.
- An explicit decision on portfolio shape — how many transformational redesigns, how many efficiency plays, how many enabling investments.
- A clear list of what is being deprioritised — with the rationale documented.
- Named senior owners for the prioritised initiatives — not project managers, executive sponsors.
- A 90-day, 12-month, and three-year roadmap — with the courage to commit to all three horizons.

WHAT GOOD LOOKS LIKE

- The portfolio shape is deliberate, not the residue of which use cases happened to score well.
- Transformational initiatives have explicit protection — they are not allowed to be starved by efficiency plays.
- The prioritisation is revisited quarterly, not annually — conditions change, and the portfolio should reflect that.
- The programme team can articulate, in plain language, why each initiative is in the portfolio — and why others are not.
- Initiatives that fail to progress are killed cleanly, not allowed to drift.

WHERE ORGANISATIONS GO WRONG

- Saying yes to too many initiatives — producing a portfolio nobody can deliver.

- Quick wins crowding out transformational bets — because the quick wins are easier to defend in next month's steering committee.
- Prioritisation done in spreadsheets, not conversations — producing a ranked list that nobody actually owns.
- Avoiding the deprioritisation conversation — leaving every stakeholder believing their pet initiative is still alive.
- Treating prioritisation as a one-off — not as a continuous discipline of pruning and reweighting.

THREE QUESTIONS TO ASK

- Of the initiatives in our current portfolio, how many would still be there if we had to choose only the top five?
- Are our transformational redesigns protected from being starved by smaller efficiency plays — or are we relying on hope?
- When did we last formally kill an initiative that was in the portfolio but not progressing?

Step 3: Design & Reimagine.

If we were designing this from scratch, knowing what AI can do — what would it look like?

This is the step where reconceptualisation either happens or does not. It is also the step most organisations rush through, because it is uncomfortable and the work feels abstract. The pressure to move to building is intense, particularly from sponsors who measure progress in delivered systems. Resisting that pressure — protecting the time and authority to do real redesign work — is the single most important act of leadership in a hyperautomation programme.

Real redesign is structured. It begins with the outcome the organisation actually wants — not the process it currently has — and works backwards to a design that delivers that outcome with AI as a fundamental component, not as a layer added later. It involves the people whose work will change. It produces a target-state design that is materially different from the current state, not a marginally improved version of it. And it makes explicit, documented choices about which steps in the current process will not exist in the future state — because those decisions are the ones the programme will spend the next twelve months defending against the gravitational pull of the existing way of working.

WHAT DESIGN REQUIRES

- A clear articulation of the outcome the redesigned process is built to deliver — in business terms, not technology terms.
- A clean-sheet redesign exercise — not constrained by the current process at the start.
- Active involvement of the people whose work will change — not just their managers.
- Explicit decisions about steps to be removed, roles to be changed, metrics to be replaced.
- A target-state blueprint that is detailed enough to be built from — not a slide deck.
- A documented gap analysis between current and target state, with the magnitude of change honestly assessed.

WHAT GOOD LOOKS LIKE

- The target-state design is materially different from the current state — not a slightly faster version of it.
- The design includes new roles, new metrics, new customer/employee experiences — not just new technology.
- The design has been pressure-tested with stakeholders who will resist it — and survived.
- Steps that will not exist in the future state are documented — along with the rationale, so they don't reappear during build.
- The design is owned by an executive who has the authority to defend it through delivery.
- Change management is built into the design from the start, not added afterwards.

WHERE ORGANISATIONS GO WRONG

- Redesign delegated to the team that owns the current process — producing modest changes that protect the status quo.
- Skipping the clean-sheet exercise — starting redesign from the current state and producing minor modifications.
- Redesign without the people whose work changes — producing designs that fail in adoption.
- Treating the design as fixed — not allowing it to evolve as the build phase reveals new information.
- Letting the technology choice constrain the design — producing a redesign limited by what the chosen tools happen to do.
- Insufficient time on this step — most underdelivered programmes had under-scoped design phases.

THREE QUESTIONS TO ASK

- Could we describe our target-state process in a way that someone outside the organisation would recognise as fundamentally different from the current state?
- Who, by name, is accountable for protecting the design from gradual erosion during build?
- What steps in the current process have we explicitly decided will not exist in the target state — and is that decision documented somewhere we can refer to in twelve months?

Step 4: Build & Integrate.

Bringing the redesign to life across systems, data, and people.

Build is the step that gets most of the attention in transformation programmes — because it is the step that produces visible artefacts, runs the largest budgets, and involves the most people. It is also, paradoxically, the step where leadership matters least and discipline matters most. If discovery, prioritisation, and design have been done well, build is largely an execution problem with well-understood patterns. If they have not been done well, no amount of build excellence will rescue the programme.

The defining discipline of the build phase is integration. Hyperautomation is, by definition, the orchestration of multiple technologies and multiple processes as a single system. The integration work — between AI models and existing systems, between automated steps and human judgement, between data sources and decision points — is the work that determines whether the programme delivers an integrated capability or a collection of disconnected components. Most underperforming programmes had perfectly adequate components and inadequate integration.

WHAT BUILD REQUIRES

- A delivery approach matched to the redesign ambition — agile for evolving work, structured for stable work, not the other way around.
- Cross-functional teams with named owners for each component and for the integrations between them.
- Data pipelines built first, not retro-fitted — most build phases discover their data problems too late.
- Integration architecture documented up front — not improvised as components arrive.
- Iterative delivery with regular checkpoints against the target-state design — to catch drift early.
- Active change management running in parallel — not sequenced after the technical build.

WHAT GOOD LOOKS LIKE

- Components are built to a documented integration specification — not improvised.
- Data quality issues are surfaced and resolved as they arise, not allowed to accumulate.
- The build team includes operations representatives, not just technologists — ensuring the build matches operational reality.
- Decisions made during build that change the target-state design are documented — not silently absorbed.
- Integration is treated as first-class work, with named owners and explicit timelines.
- The system delivered at the end of build matches the design produced at the end of step 3 — with deviations documented and justified.

WHERE ORGANISATIONS GO WRONG

- Build starts before design is finalised — producing a system that solves the wrong problem competently.
- Components built independently and integrated last — producing integration nightmares that consume the final third of the timeline.
- Data work deferred — discovered as a programme-stopping issue late in delivery.
- Change management sequenced after build — producing technically successful systems that fail in adoption.
- Vendor-led builds without strong client-side architecture authority — producing systems shaped to vendor capabilities rather than business needs.
- Documentation deprioritised in favour of speed — producing systems nobody can maintain, audit, or improve.

THREE QUESTIONS TO ASK

- Is the build phase delivering against the target-state design — or is the design being quietly adapted to match what the build is producing?
- How is integration being treated in the programme — as a discipline with named owners, or as a problem we will solve at the end?
- Are the people whose work will change involved in the build — or will they meet the system for the first time at deployment?

Step 5: Test & Validate.

Does the redesigned process actually deliver what the design promised?

Test and validation are the disciplines that separate competent transformation programmes from negligent ones. Most programmes test the technology — does the system work, does it produce the right outputs, does it integrate with the surrounding stack. Far fewer programmes test the redesign — does the new process actually deliver the business outcomes the design promised, across all the conditions it will encounter, across all the stakeholders it affects. Testing the technology and not testing the redesign is the most common reason that successful go-lives produce disappointing results.

Real validation also tests the things that are easy to miss: bias and fairness in AI components, edge cases that fall outside the training data, security and adversarial robustness, the human handover points where AI hands work to people and vice versa. Each of these is its own discipline, and each can be the source of a failure mode that an outputs-only test would miss.

WHAT TESTING REQUIRES

- Defined acceptance criteria, written before testing begins — not negotiated based on what the system happens to do.
- Testing across the full range of conditions the system will encounter — not just the easy cases.
- Bias and fairness testing for AI components — across the populations the system will affect.
- User acceptance testing with the actual users, not their managers — in their actual environment, not a sanitised one.
- End-to-end process validation — testing the whole redesigned workflow, not just its components.
- A documented go/no-go decision based on evidence — with the evidence preserved.

WHAT GOOD LOOKS LIKE

- The system meets or exceeds the success criteria defined in steps 1 and 3 — measured, not asserted.
- Edge cases and failure modes are documented — not discovered after deployment.
- Bias has been tested for and addressed — with evidence retained for future review.
- Operational readiness is verified, not assumed — the people, the process, and the technology all ready together.
- The decision to deploy is defensible — if questioned in twelve months, the rationale and evidence can be produced.

WHERE ORGANISATIONS GO WRONG

- Testing only what is easy to test — producing systems that pass technical tests and fail in operational reality.

- Acceptance criteria written after testing — producing standards the system happens to meet.
- Skipping bias testing — because it is harder, less standardised, or politically inconvenient.
- User acceptance testing with proxies, not real users — producing approval that does not predict adoption.
- Treating test failures as obstacles to deployment, not data about the system — producing pressure to soften the criteria.
- No documented go/no-go decision — producing programmes that drift into deployment.

THREE QUESTIONS TO ASK

- What did we deliberately decide not to test — and what risk are we accepting by that decision?
- Have we tested the redesigned process end-to-end with real users in real conditions — or only its components in controlled environments?
- If this system causes harm in the first month of operation, what evidence will we have that we tested for it?

Step 6: Deploy & Scale.

Moving from validated capability to live organisational reality.

Deployment is the moment the programme stops being a project and becomes operational reality. It is also the moment when the discipline of the previous five steps either holds up or unravels. A programme that has done good discovery, real prioritisation, honest design, careful build, and rigorous test can still fail in deployment if the operational handover is botched — if oversight is not in place, if the people running the new process are not ready, if monitoring is not active, if the conditions for rollback are not defined.

Scale is the work that comes after first deployment. Most successful pilots fail to scale because the work to scale is qualitatively different from the work to pilot. Pilots reward narrow scope, dedicated teams, and tolerance for hand-holding. Scale rewards standardisation, broad coordination, and friction tolerance. Programmes that are deployed without an explicit scale strategy tend to remain pilots — successful in a small corner of the organisation, never reaching the size where they would justify the investment.

WHAT DEPLOYMENT AND SCALE REQUIRE

- A controlled go-live, with named accountability for the operational system from day one.
- Monitoring active from the moment of deployment — not added later.
- Defined rollback conditions and a tested rollback procedure.
- Operational readiness verified — the people, the process, the supporting systems all aligned.
- A separate scale plan, with its own owner and its own timeline — not a footnote on the deployment plan.
- Communication to all affected stakeholders — customers, employees, partners — calibrated to each audience.

WHAT GOOD LOOKS LIKE

- The deployed system is owned, monitored, and supported from day one — not handed over to operations as a surprise.
- Performance against the success criteria from step 1 is being measured and reported, with action triggers if metrics deteriorate.
- Scale follows a documented plan — not opportunism.
- Lessons from first deployment are captured and applied to subsequent deployments — systematically.
- The redesigned process is the dominant way of working within twelve months of deployment — not an option people can opt out of.

WHERE ORGANISATIONS GO WRONG

- Deploy and forget — the team disbands, monitoring is not maintained, the system drifts.

- Scale by accident — each new deployment improvises rather than learning from the prior ones.
- Permitting the old process to run alongside the new one — indefinitely — producing two ways of working and the value of neither.
- No defined point at which the programme ‘ends’ — producing perpetual transformation rather than embedded change.
- Insufficient communication to customers and employees — producing avoidable trust failures.
- No reflection between deployments — reproducing the same problems at scale that were visible at pilot.

THREE QUESTIONS TO ASK

- On day one of deployment, who answers the phone if the new system fails — and is that named, by person?
- What is our scale plan, and is it owned by someone whose performance is measured against scale outcomes — or by the team that built the pilot?
- When does the old process get switched off — and who has the authority to switch it off?

Step 7: Optimise & Evolve.

Hyperautomation is a discipline, not a project. The seventh step never ends.

The seventh step is the one that turns a successful transformation into a sustained capability. Without it, even excellent programmes deliver one-off improvements that gradually decay. With it, the organisation builds a compounding asset — a redesigned process that continues to learn, evolve, and improve, producing value that grows rather than flattens over time.

The discipline of optimise and evolve has three components. The first is operational monitoring — watching the live system for drift, degradation, and emerging risk, with thresholds that trigger action. The second is continuous improvement — systematically harvesting the lessons of operation and feeding them back into the design, the build, the deployment of subsequent capabilities. The third is horizon scanning — watching the technology and regulatory environment for changes that affect the design choices, and updating the system before those changes become problems. Programmes that do all three produce compounding value. Programmes that do one or two produce diminishing value. Programmes that do none produce decaying value.

WHAT OPTIMISE AND EVOLVE REQUIRE

- Continuous monitoring of system performance against the original success criteria — not a one-off post-implementation review.
- A documented continuous improvement process — with named owners and a regular cadence.
- Lessons learned from incidents and near-misses captured and applied — not anecdotally remembered.
- Regular re-validation of the design as conditions change — not a fix-and-forget posture.
- Horizon scanning for technology and regulatory developments — with a process to absorb them.
- Connection back to discovery — each operational system surfaces new opportunities that feed the next round of the playbook.

WHAT GOOD LOOKS LIKE

- Performance against original success criteria is reported regularly to executive level — and acted on when it deteriorates.
- Improvement actions are documented, sequenced, and visible — not informal.
- Each operational system feeds insight back into the discovery process — producing a continuously refreshed portfolio.
- The organisation has a hyperautomation capability that has outlived its original sponsor — institutionalised, not personal.
- Compounding value is visible in the metrics — each year better than the last, not flat.

WHERE ORGANISATIONS GO WRONG

- Treating deployment as the end — disbanding the team, ending the funding, declaring victory.
- Monitoring exists but isn't acted on — producing dashboards that nobody reads.
- Improvement happens reactively, when something breaks — not proactively, as discipline.
- No connection back to discovery — each new initiative starts from scratch instead of from accumulated learning.
- Letting the operational system drift toward its original shape — because nobody owns its continued evolution.
- Confusing absence of incidents with presence of improvement — the system that doesn't fail visibly but no longer learns is decaying.

THREE QUESTIONS TO ASK

- For our most material redesigned process, what is the cadence of formal review — and who attends it?
- Of the lessons from incidents and near-misses in the past year, how many have produced documented changes to the system?
- Could we name three improvements made to a deployed AI system in the past quarter — and the people responsible for each?

> *"The first deployment proves the system works. The seventh step is what makes the value compound."*

CHAPTER 5

The foundational enablers.

Five conditions that determine whether the playbook delivers value or stalls.

The seven steps describe what to do. The five enablers below describe the conditions that have to be in place for the work to succeed. Each is a programme in its own right, with its own owner and its own investment case. Each is also commonly underestimated, and the underestimation is one of the most consistent reasons hyperautomation programmes deliver less than they promised. Treat the enablers as parallel investments, not as a checklist to clear before starting.

Enabler 1: Data.

Data is the rate-limiting factor on every hyperautomation programme. AI models trained on poor data produce poor outputs. AI systems integrated with siloed data produce siloed outcomes. Process redesigns that depend on data the organisation does not have, or cannot access, or is not allowed to use, fail in ways that look like technology problems but are actually data problems. The organisations that get hyperautomation right invest disproportionately in their data foundation, and they do so before the AI investment, not in parallel with it.

This is not a call for a multi-year data warehouse programme. It is a call for clear-eyed honesty about which data assets are genuinely AI-ready — accurate, accessible, current, governed — and which are not, and what it would take to get the critical ones to ready. Hyperautomation programmes that gloss over this assessment discover its truth in the build phase, when the data they assumed would be available proves to be inaccessible, dirty, or legally constrained. By that point, the programme is committed to a target state it cannot reach.

Enabler 2: Technology.

The technology layer for hyperautomation is more complex than the layer for traditional automation, because hyperautomation is by definition the orchestration of multiple technologies as a coordinated system. AI and machine learning, robotic process automation, business process management, integration platforms, low-code tools, intelligent document processing, analytics and observability — these are not alternatives to each other, they are components that need to work together. The capability the organisation needs is not just any one of these but the architectural discipline to make them combine into a coherent platform.

This means the technology question is, ultimately, an architecture question. The right answer is rarely ‘buy the suite from a single vendor’ — because no single vendor is best at all components, and single-vendor lock-in produces its own forms of fragility. The right answer is also rarely ‘let each team choose its own tools’ — because that produces a stack that cannot be orchestrated. Mature hyperautomation organisations have an explicit architectural pattern, with named components, defined integration standards, and authority to enforce them across teams.

Enabler 3: People and culture.

Hyperautomation changes work. The people whose work it changes are the determining factor in whether the change actually delivers value or whether it produces a system that technically works and operationally fails. Investment in the people side of transformation is the second most-common area of underinvestment after data, and the consequences are the same: programmes that look good on paper and underperform in operation.

The people investment has three components. Capability — the skills required to operate the redesigned process, including AI fluency at the right level for each role. Adoption — the change management work required to move from the current way of working to the new one. Culture — the broader organisational disposition toward innovation, experimentation, and continuous improvement. Each component requires its own investment, and substituting one for another does not work. A programme that runs adoption well but fails on capability produces frustrated users; a programme that builds capability but fails on culture produces individuals who improve while the organisation does not.

Enabler 4: Governance.

Governance is the enabler most often confused for an enemy of transformation, particularly in early stages of a programme when momentum feels precarious. The instinct is to defer governance until after the value has been demonstrated, on the theory that governance will slow things down. The instinct is wrong. Governance designed in from the start is significantly cheaper, significantly more effective, and significantly less obstructive than governance retro-fitted to operational systems. Governance retro-fitted is also significantly more visible to regulators, auditors, and customers — because it shows.

Hyperautomation programmes designed without governance produce, on average, three predictable failure modes: inability to demonstrate compliance to regulators or customers; inability to investigate or explain individual AI decisions; and inability to retire systems cleanly when they reach the end of their useful life. Each is expensive to fix. The fix involves doing the governance work that was deferred, plus paying for the operational consequences of having deferred it. The honest cost of ‘we’ll add governance later’ is several multiples of the cost of designing it in. The Executive’s Guide to AI Governance, the first volume in our parallel Executive Series, sets out the governance framework that complements this playbook.

Enabler 5: Partnerships.

Hyperautomation is rarely done alone. The technology landscape is complex enough that few organisations have all the capabilities in-house, and even those that do benefit from external perspectives that an internal team cannot provide. The fifth enabler is the discipline of choosing partners well, contracting with them appropriately, and managing the relationship with the same rigour as any other strategic dependency.

Strong partnerships in hyperautomation share three characteristics. The partner is paid for outcomes, not activity. The partner's success criteria are aligned with the organisation's success criteria — not with delivering more of the partner's product. And the partner is contributing capability the organisation could not produce internally — not duplicating capability the organisation could build for itself. Partnerships that fail these tests produce predictable patterns of disappointment, and the disappointments tend to surface late, after switching costs have accumulated.

THE FIVE ENABLERS — A BOARD CHECK

- Data — do we have an honest assessment of what is AI-ready, and an investment plan for what is not?
- Technology — do we have an explicit architecture, with authority to enforce it?
- People and culture — are capability, adoption, and culture each invested in, or are we relying on one to substitute for the others?
- Governance — is governance designed in from the start, or are we deferring it to after value is demonstrated?
- Partnerships — are our partners contracted for outcomes, aligned to our success, and contributing genuine capability we could not produce internally?

If the answer to any of these is uncomfortable, the enabler is underinvested. The good news is that all five are knowable and addressable. The work is to address them deliberately, before the programme depends on them.

CHAPTER 6

Before and after: a worked example.

What the difference between automation and hyperautomation actually looks like.

The purpose of this chapter is to illustrate the difference that the rest of the playbook describes in principle. The example below is a generalised business process — customer onboarding for a mid-sized financial services or professional services firm — chosen because the pattern is recognisable across many industries. The specifics will not match any particular organisation. The pattern will match many.

The starting point.

In the before state, customer onboarding is a sequence of handoffs. The customer fills in an application form, online or on paper. The form arrives in a queue. A staff member reviews it for completeness and follows up on missing information. Once complete, identity verification happens through a combination of automated and manual checks. A risk assessment is generated, drawing on data from internal systems and external bureaux. The case is reviewed by a manager. Compliance does a separate review. The customer is notified. End-to-end, the process takes between three and fifteen working days, with significant variation depending on case complexity. Each step is documented. Each step is necessary, given how the others work. None of the steps, individually, is unreasonable.

The cost of this process is rarely calculated end-to-end, but when it is, the numbers are uncomfortable. Staff time across the handoffs. Customer drop-off during the wait — a meaningful percentage of customers who complete the application never complete the onboarding, because the wait exceeds their patience. Operational risk from manual handoffs. Compliance overhead. Capacity constraints during peak periods. The process delivers a customer onboarded, but at a cost in money, time, and customer experience that an outside observer would describe as substantial.

The automation answer.

The automation approach to this process is straightforward and limited. AI is applied to specific steps. Form completion is improved with intelligent autofill and validation. Identity verification is automated where possible. Risk scoring uses an AI model trained on historical data. Document processing extracts information from uploaded files automatically. The manual review step remains, but the reviewer now sees AI-generated summaries and recommendations. The compliance step remains, but with similar AI-assisted summaries. Customer notifications are drafted by AI and sent automatically once approved.

The result is a faster version of the same process. End-to-end time drops from three-to-fifteen days to one-to-five days. Staff time per case falls by perhaps thirty per cent. Customer drop-off improves modestly. The dashboard metrics look good. The board approves the next phase. The cost per customer onboarded has fallen, but the structure of the experience — sequential review, multiple touchpoints, days rather than minutes — is unchanged. The customer still applies, waits, and is notified later. The competitive position of the organisation, against a peer that has reconceptualised the process, is unchanged or worse.

The hyperautomation answer.

The hyperautomation answer is a different process. The customer is onboarded in the moment of application. AI handles identification, validates documentation, scores risk, and reaches a decision in real time. Where the AI is confident, the decision is final and the customer is onboarded immediately. Where the AI is uncertain, the case is routed to a human reviewer with a structured summary of the uncertainty and the data needed to resolve it. The human reviewer's role has changed: from processing every case to handling the meaningful exceptions. Compliance is no longer a separate step; it is built into the decisioning logic, with full audit trail produced as a by-product. The customer experience is not 'faster onboarding'. It is 'onboarded now, like any modern digital service'.

The numbers reflect the difference. End-to-end time for routine cases is minutes, not days. Staff time per case falls by seventy to ninety per cent for routine cases, and by less for exceptions — but the exceptions are a small fraction of total volume. Customer drop-off falls dramatically because the waiting that caused it no longer exists. Customer experience is competitive with digital-native peers, not behind them. The cost structure of the function is fundamentally different: a smaller team handling a larger volume, focused on exceptions and improvements rather than routine processing.

Most importantly, the redesigned process produces something the automated process does not: continuous learning. Every routine case is data that improves the model. Every exception is a signal about where the model is uncertain, which is a signal about where to invest in further capability. The process is not just operating; it is improving, year on year, while the automated alternative quietly decays as conditions change.

WHAT CHANGED BETWEEN THE TWO ANSWERS

- Steps removed — not just sped up. The wait, the queue, the sequential reviews are not faster. They no longer exist.
- Roles changed. Reviewers handle exceptions, not routine cases. Their work is more complex, less repetitive.
- Metrics changed. The KPIs are not 'cycle time' anymore. They are 'instant onboarding rate', 'exception rate', 'model confidence'.

- Customer experience changed in kind, not in degree. Onboarding is now an event, not a process.
- Organisational capability changed. The team has evolved from operators to improvers.
- Continuous improvement is built in. Every case is data. Every exception is a signal. The system gets better with use.

This is what the McKinsey finding cited in the foreword measures. Workflow redesign is not faster steps. It is fewer steps, different steps, different roles, different metrics, different customer experience. The organisations that do it produce returns three times more often than the organisations that do not. The example above generalises. The pattern is what matters.

> *"Automation gets you to the same destination faster. Hyperautomation gets you to a different destination."*

CHAPTER 7

The governance and fluency layer.

Transformation without governance is risk. Transformation without fluency is failure. Both have to be present.

This is the shortest chapter in the playbook, and the most important to get right. Hyperautomation, done well, produces transformational value. It also produces transformational risk: AI systems making consequential decisions at scale, on behalf of the organisation, in ways that affect customers, employees, regulators, and stakeholders. The discipline that contains that risk is governance. The discipline that ensures the humans inside the redesigned processes can operate them well is fluency. Both are subjects in their own right, covered in detail in our parallel Executive Series. This chapter is the bridge.

The governance dimension.

Hyperautomation programmes that ignore governance produce systems that cannot survive their first material incident. The transformations work; the governance does not. When the inevitable question arrives — from a regulator, an auditor, a major customer, an employee whose work has changed in ways they did not consent to — the organisation cannot answer it. The redesigned processes deliver real value, but they deliver it in a way that the organisation cannot defend. That is not transformation. It is borrowing against future trust to fund current performance.

Governance designed into the hyperautomation programme from the start — not retro-fitted after deployment — produces three outcomes. First, the redesigned processes are auditable, explainable, and defensible. Second, the regulatory environment that is coming — ISO/IEC 42001 in 2027 to 2028 in South Africa, the EU AI Act already in force, sector-specific regulation accumulating — is something the organisation is ready for, not surprised by. Third, the trust the redesigned processes earn from customers, employees, and stakeholders is durable rather than precarious. The first volume of our Executive Series, The Executive's Guide to AI Governance, sets out the ten-pillar framework that supports this work.

The fluency dimension.

Hyperautomation programmes that ignore fluency produce systems that work in principle and fail in practice. The redesigned processes depend on humans — to handle exceptions, to oversee outputs, to intervene when something goes wrong, to improve the system over time. Those humans need fluency: the practical capability to work with AI competently, ethically, and safely. The four-competency AI fluency framework — Delegation, Description, Discernment, Diligence — and applied to executive contexts in our Executive's Guide to AI Fluency for Boards, sets out what that fluency consists of: Delegation, Description, Discernment, Diligence.

Fluency in a hyperautomation programme is most acute at three points. The discovery and design stages, where Delegation — the discipline of choosing what AI should and should not do — determines the shape of the redesign. The operation stage, where Discernment — the ability to recognise when AI outputs are wrong, biased, or unfit — protects the organisation from confidence-shaped errors. And the continuous improvement stage, where all four competencies cycle through every redesigned process, every day, in the hands of the people who run it. Programmes that build fluency deliberately produce these capabilities by design. Programmes that do not get them by accident, when they get them at all.

HOW THIS PLAYBOOK INTEGRATES WITH THE EXECUTIVE SERIES

- The Executive's Guide to AI Governance sets out the ten-pillar framework for governing AI — the controls that this playbook's redesigned processes must operate within.
- The Executive's Guide to the AI Product Lifecycle sets out the seven-stage discipline for governing AI systems through their life — the operating manual for the systems this playbook produces.
- The Executive's Guide to AI Fluency for Boards sets out the human capability layer — the four competencies that the people inside this playbook's redesigned processes need to operate them well.
- Together with this playbook, the four guides cover what to govern, how to govern it, who has to be capable of governing it, and how to deliver transformational value — the four sides of a complete AI strategy.

Read alone, this playbook will help you deliver real transformation. Read alongside the Executive Series, it will help you deliver real transformation that survives contact with regulators, customers, and time. Both readings are valid. Which is appropriate depends on the organisation, the stakes, and the maturity of the existing governance environment. For most mid-market organisations in South Africa today, both readings are necessary.

CLOSING

Three commitments for transformation leaders.

If the rest of this playbook is correct, three things follow. Each is the personal commitment of the leader who sponsors the work.

Hyperautomation is not a programme management problem. It is a leadership problem. The technical disciplines are well-established and broadly available. The frameworks are settled. The tools are mature. What is missing in the organisations that underdeliver — and present in the organisations that produce transformational results — is the leadership posture this closing chapter describes. Three commitments. Each is personal. Each is testable.

Commitment one: redesign before you automate.

The first commitment is the central argument of this playbook: that the value from AI is in the redesign, not in the automation, and that any programme which does not protect the redesign work will deliver a slightly faster broken workflow. The leader who makes this commitment refuses to fund tool-first programmes. Refuses to accept use cases that do not include reconceptualisation. Protects the design phase from the gravitational pull of ‘let’s just automate what we have for now’. Holds the line when the pressure to ship overrides the pressure to redesign. This commitment is mostly about saying no — to programmes, vendors, and stakeholders that want to skip the harder work.

Commitment two: invest in the conditions, not just the programme.

The second commitment is to fund the five enablers as parallel investments, not as a checklist to clear before starting. Data foundations, technology architecture, people and culture, governance, partnerships — each is an investment in its own right, and underinvesting in any of them places a ceiling on what the programme can deliver. The leader who makes this commitment treats the enablers as part of the transformation budget, not as overheads against it. Defends them at executive level when budget pressure tries to compress them. And measures progress on them as rigorously as progress on the headline programme.

Commitment three: treat hyperautomation as a permanent capability.

The third commitment is the one that turns successful programmes into compounding ones. Hyperautomation is not a project that finishes. It is a permanent organisational capability, with its own funding, its own talent, its own governance, its own evolution. The leader who makes this commitment refuses to disband the team after first deployment. Insists on continuous improvement infrastructure as part of the design. Builds an institution that will outlive the original sponsor and continue to deliver value years after the founding programme has finished. This is the commitment that turns a transformation success into a transformation advantage.

> *“The technical disciplines are settled. The frameworks are mature. What separates the organisations that produce transformational results from the organisations that produce slightly faster broken workflows is leadership posture — nothing more complicated than that.”*

This is the first volume of The Transformation Series. It is designed to be read alongside the Executive Series, but it can be read alone by leaders whose remit is value rather than control. The next volume in the series will go deeper into one specific dimension of the work — likely process redesign as a discipline — with the same standard of rigour as this one.

Until then, the work is in front of every transformation leader who has read this far. The playbook is here. The thesis is sharp and the frameworks are mature. The choice between automation and hyperautomation is a choice the organisation makes — deliberately or by accident, every time it scopes a programme. The leaders who choose deliberately are the ones who will produce transformational results in the years ahead. The rest will produce, on average, slightly faster broken workflows. The McKinsey numbers say so.

DILIGENCE STATEMENT :

In putting this playbook together, I made use of Claude to augment this artifact, Chat GPT created the images. I take full responsibility for the output of this augmentation and undertake to ensure that it is regularly updated.

Where are you now? A seven-step self-assessment.

For each step in the seven-step playbook, mark your organisation's current position honestly. This is a map, not a test. The steps marked Not in place or Partially in place are where to focus first.

Programme Step	In place	Partially in place	Not in place
Step 1: Discover & Assess	☐	☐	☐
Step 2: Prioritise & Select	☐	☐	☐
Step 3: Design & Reimagine	☐	☐	☐
Step 4: Build & Integrate	☐	☐	☐
Step 5: Test & Validate	☐	☐	☐
Step 6: Deploy & Scale	☐	☐	☐
Step 7: Optimise & Evolve	☐	☐	☐

0–2 steps In place: begin with Discovery and Design before anything else — these are the foundations the other steps depend on. 3–5 steps In place: Design and Build discipline is typically the priority gap. Protect the design phase and your continuous improvement infrastructure. 6–7 steps In place: the Optimise and Evolve step is where most programmes eventually plateau. Ask whether your redesigned processes are genuinely improving year on year — or just operating.

For a facilitated version of this assessment applied to a specific transformation opportunity, take the AI Fluency Assessment at https://zoefrancois123-netizen.github.io/AI-Fluency-Company/ or book a 30-minute conversation at https://calendly.com/zoefrancois123/30min.

ABOUT

The AI Fluency Company.

The AI Fluency Company helps executive teams adopt AI responsibly — building the language, the judgement, the governance, and the transformation capability to make AI deliver real business value. We work with mid-market organisations, family businesses, and compliance-minded leaders who want to do AI properly the first time.

How we work with executive teams.

ENGAGEMENT	WHAT IT IS
AI Fluency Briefing	For Leadership teams who know AI matters, but they do not yet share language, boundaries, decision criteria, or a practical view of what AI should and should not do in the organisation. This is a half-day workshop for leadership teams. Builds shared understanding of working with AI effectively, efficiently, ethically, and safely.
AI Governance Readiness Diagnostic	Many organisations have AI activity already underway in their businesses. Tools, vendors, pilots, staff experimentation, and automated decisions, are in place but ownership, controls, risk, and value are not being accounted for.
AI Strategy Workshop	The organisation has interest, ideas, and perhaps pilots, but lacks a disciplined way to evaluate and prioritise AI opportunities, size value, assign ownership, and decide what should move forward. This two day workshop provides methodology to prioritises the use cases, map ownership, build roadmaps.
Hyperautomation Opportunity Sprint	Teams are using AI to speed up fragments of work, but the larger workflow remains slow, manual, fragmented, or badly designed. The value is trapped in the process, not the tool. Structured engagement to apply the principles in this playbook.
AI Governance Implementation Retainer	If after the diagnostic or workshop, the organisation requires sustained help translating recommendations into policy, controls, ownership, reporting, and decision habits a retainer option is available.

About the author.

Zoë François is the founder of The AI Fluency Company. She trained as a chemical engineer and built one of South Africa's earliest applied industrial AI systems in the 1990s — an expert system on the Gensym G2 platform that modelled Anglo American Platinum's refinery operations. She then spent two further decades in commercial leadership, ultimately as **General Manager of Sales at Sappi** and **Senior Sales Executive at Consol Glass**. She holds the Oxford Saïd Business School AI Programme certificate.

Engineering rigour. Commercial judgement. The combination most AI engagements are missing.

TAKE THE NEXT STEP

Take the AI Fluency Assessment for a maturity benchmark: https://theaifluencyco.com

Book a 30-minute conversation to walk through where your transformation programme is today — and where the largest opportunities for redesign lie. No charge for the first call.

Read the companion Executive Series — The Executive's Guide to AI Governance, The Executive's Guide to the AI Product Lifecycle, and The Executive's Guide to AI Fluency for Boards — for the governance, lifecycle, and fluency disciplines that complement the work in this playbook.

First edition · 2026 · Cape Town